

Overview of Java 8 Parallel Streams

(Part 1)

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

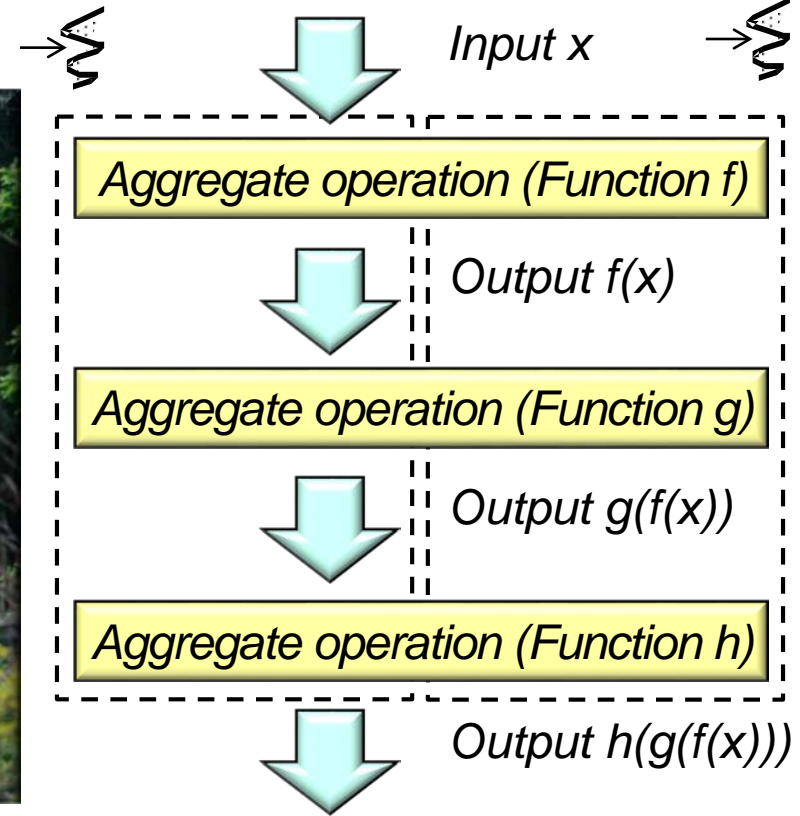
Institute for Software
Integrated Systems

Vanderbilt University
Nashville, Tennessee, USA



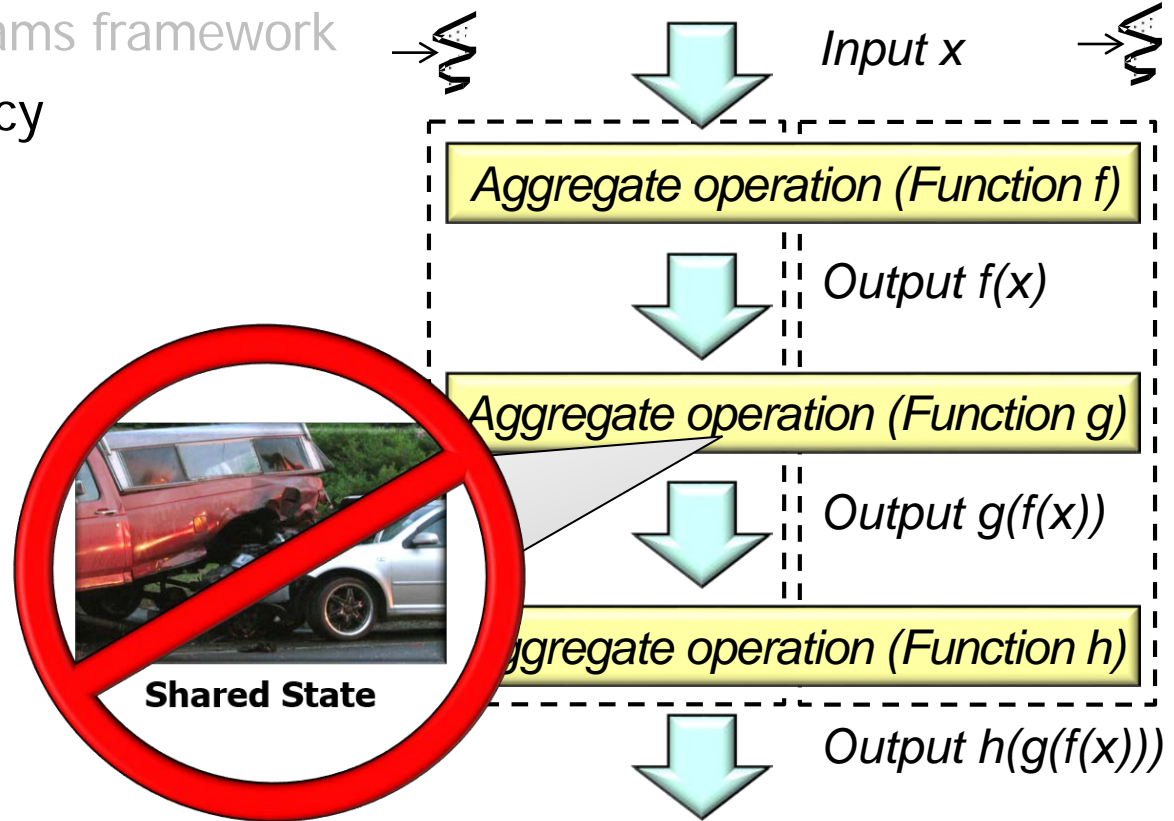
Learning Objectives in this Part of the Lesson

- Recognize how Java 8 applies aggregate operations & functional programming features in the parallel streams framework



Learning Objectives in this Part of the Lesson

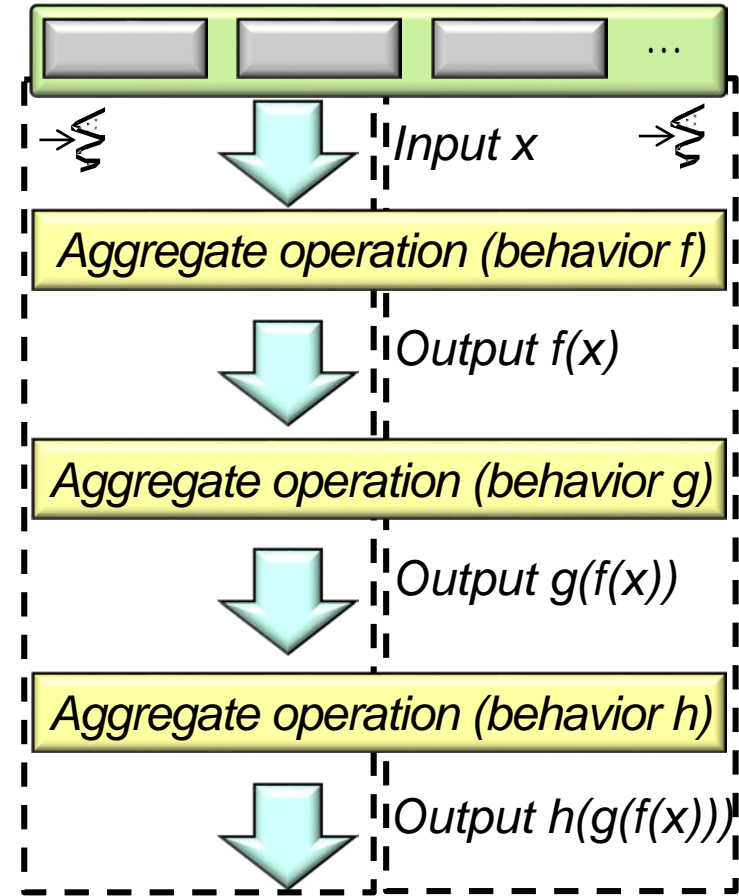
- Recognize how Java 8 applies aggregate operations & functional programming features in the parallel streams framework
- Be able to avoid concurrency hazards in parallel streams



Overview of Java 8 Parallel Streams

Overview of Java 8 Parallel Streams

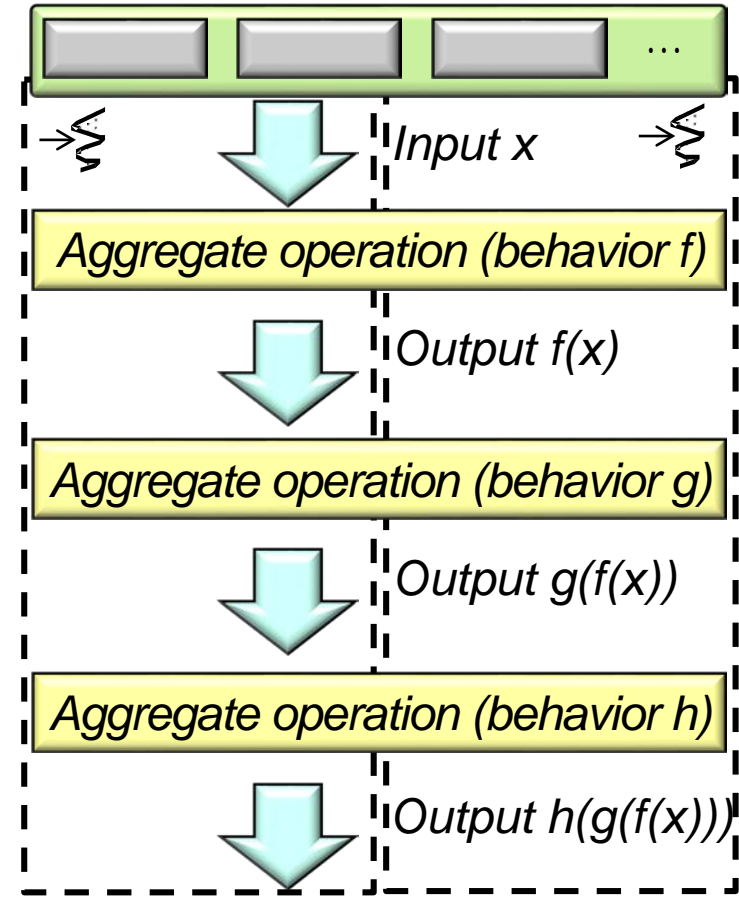
- A Java 8 parallel stream splits its elements into multiple chunks & uses a thread pool to process these chunks independently



See docs.oracle.com/javase/tutorial/collections/streams/parallelism.html

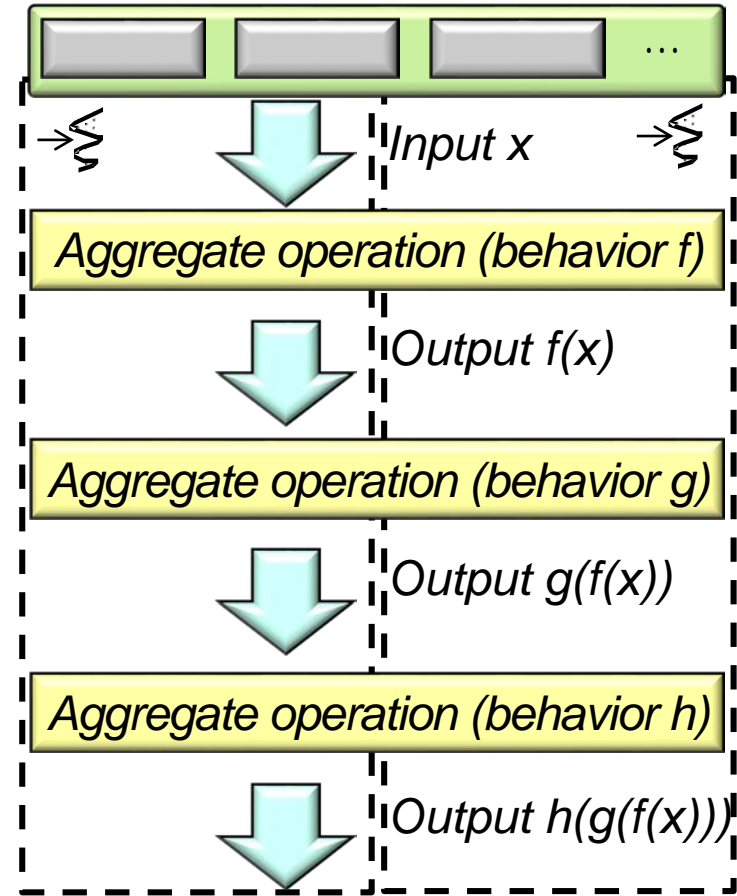
Overview of Java 8 Parallel Streams

- A Java 8 parallel stream splits its elements into multiple chunks & uses a thread pool to process these chunks independently
- This splitting & thread pool are often invisible to programmers



Overview of Java 8 Parallel Streams

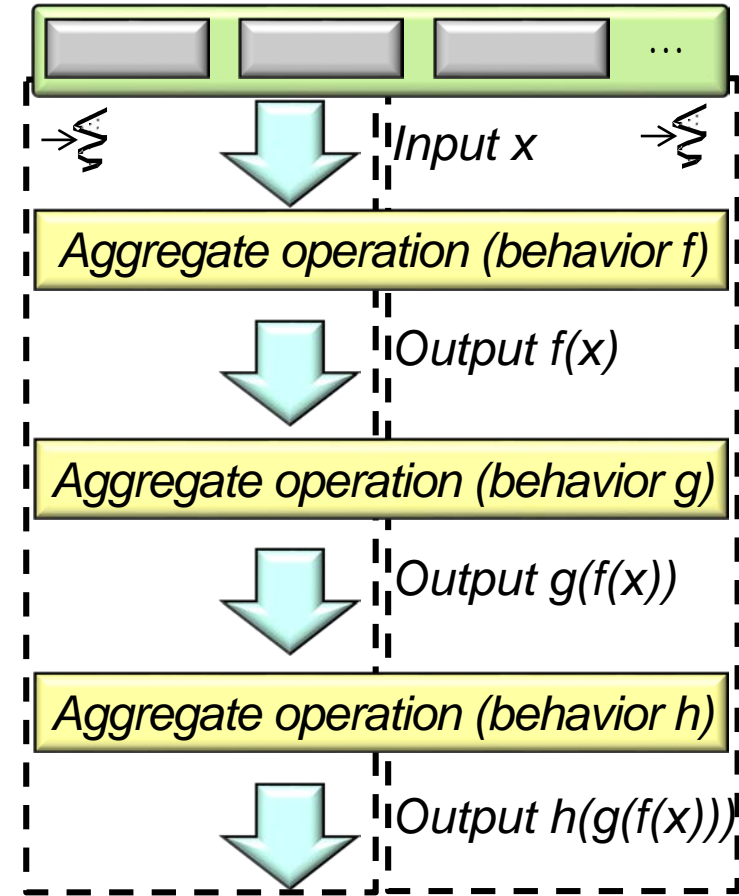
- A Java 8 parallel stream splits its elements into multiple chunks & uses a thread pool to process these chunks independently
 - This splitting & thread pool are often invisible to programmers
- The *order* in which chunks are processed is likely non-deterministic



i.e., programmers often have little/no control over how chunks are processed

Overview of Java 8 Parallel Streams

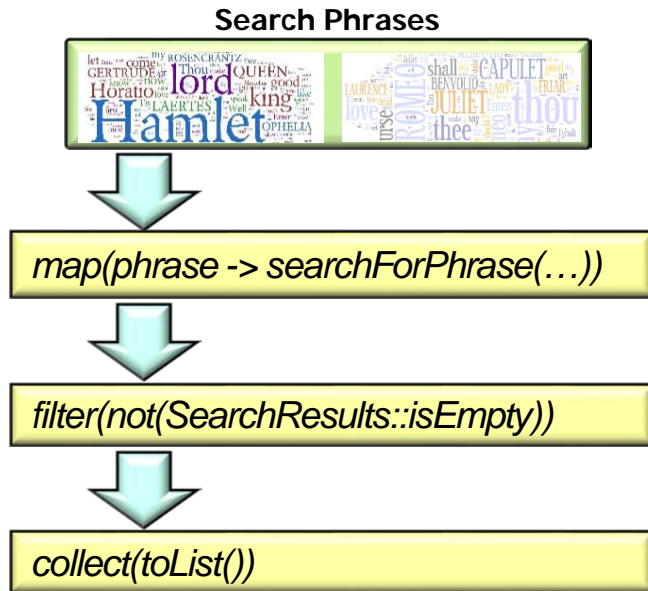
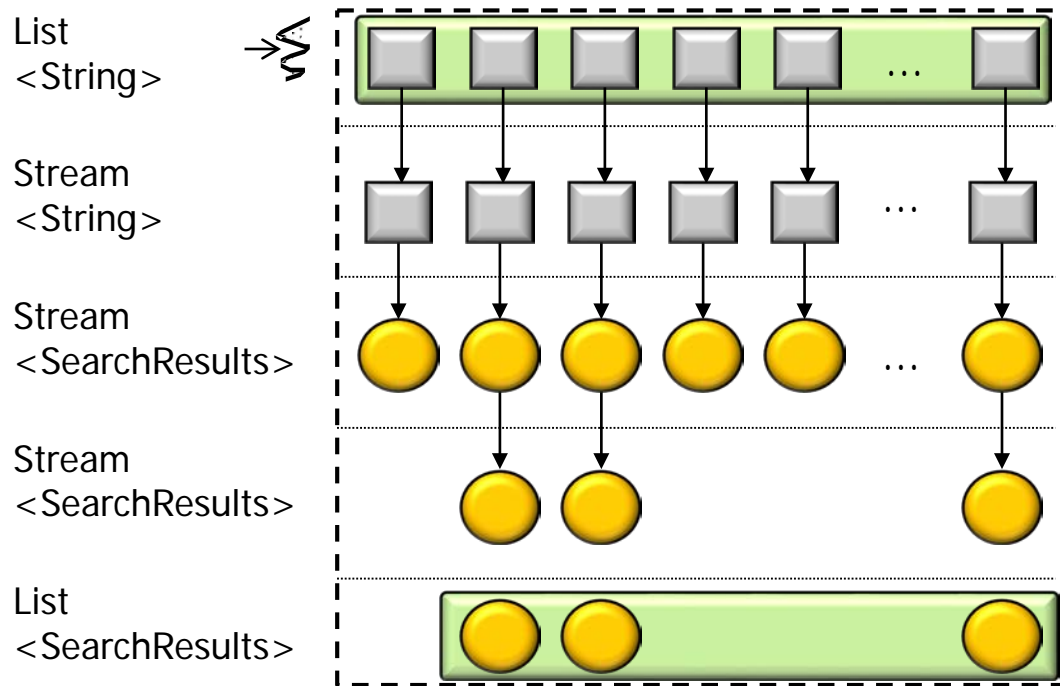
- A Java 8 parallel stream splits its elements into multiple chunks & uses a thread pool to process these chunks independently
 - This splitting & thread pool are often invisible to programmers
 - The *order* in which chunks are processed is likely non-deterministic
- The *results* of the processing is likely deterministic



Programmers have more control over how the results are presented

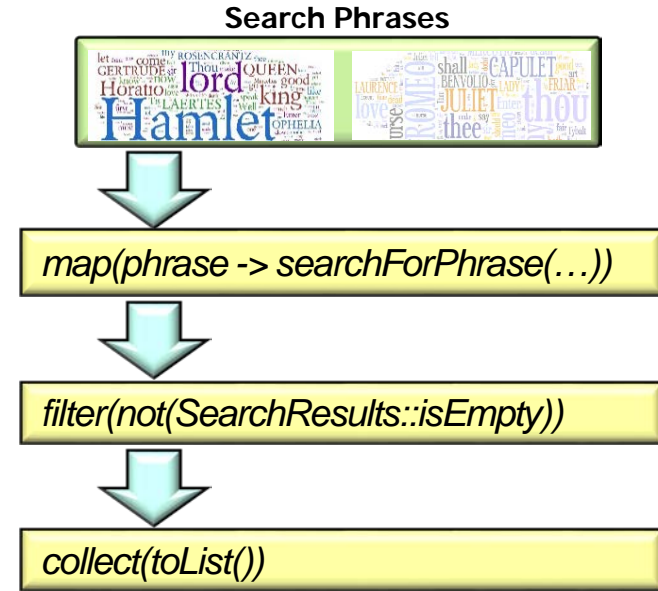
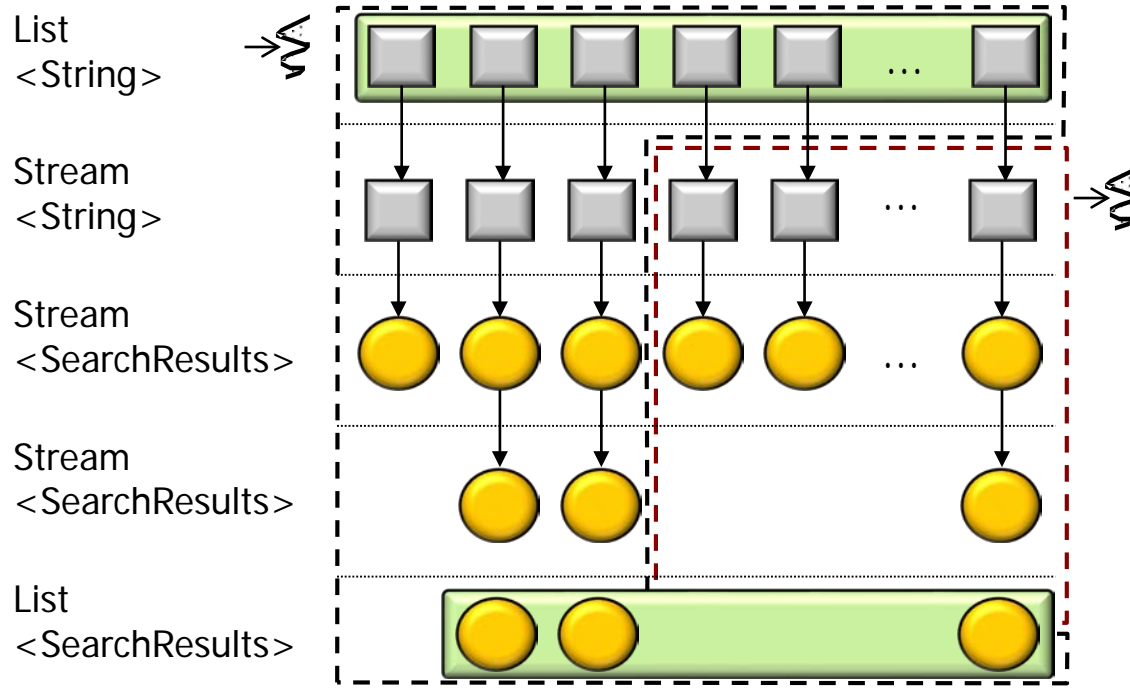
Overview of Java 8 Parallel Streams

- When a stream executes sequentially all of its aggregate operations run in a single thread



Overview of Java 8 Parallel Streams

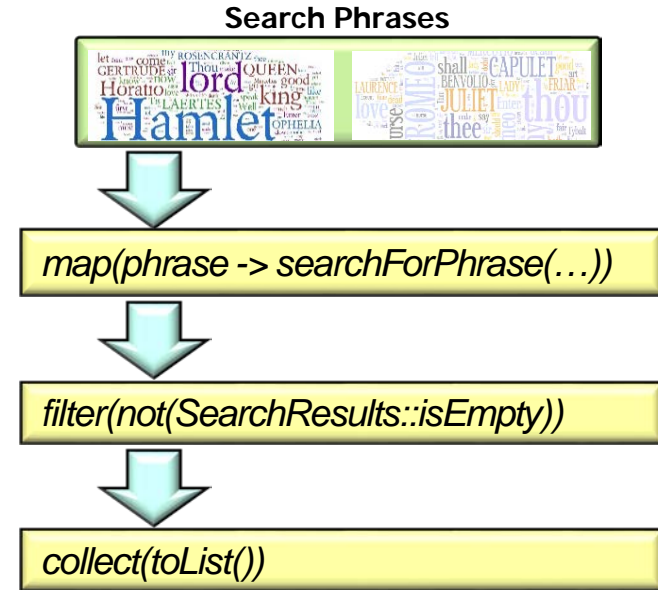
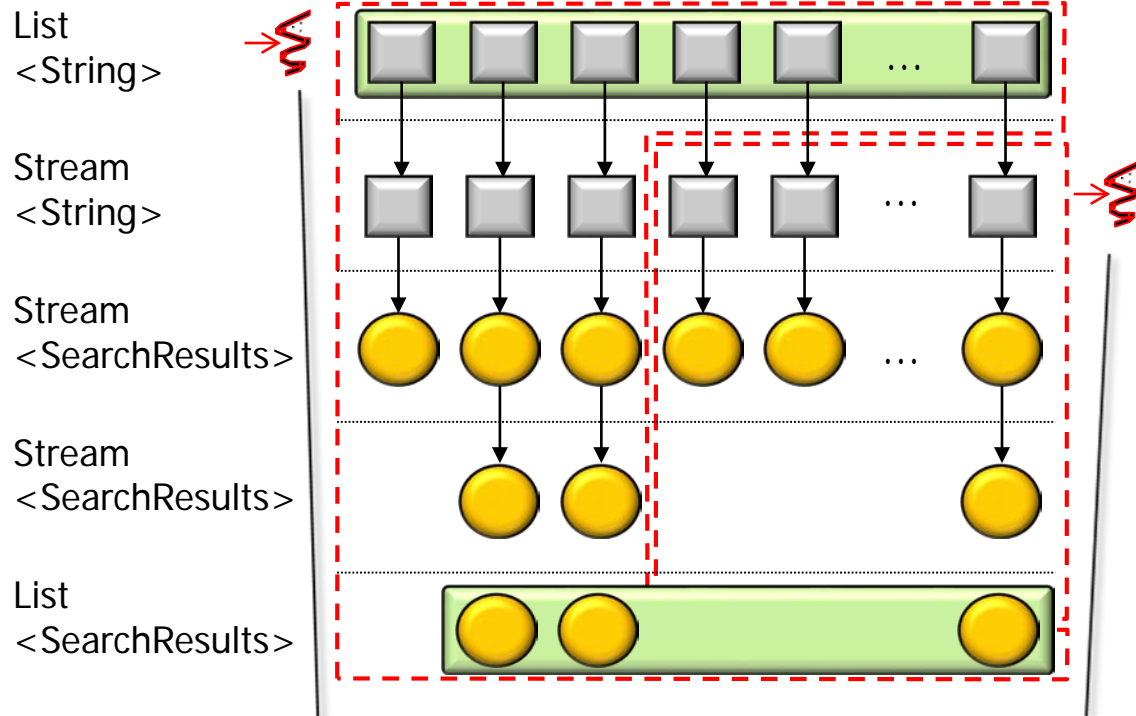
- When a stream executes in parallel, the Java runtime partitions it into multiple substream “chunks” that run in a common fork-join pool



See docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinPool.html

Overview of Java 8 Parallel Streams

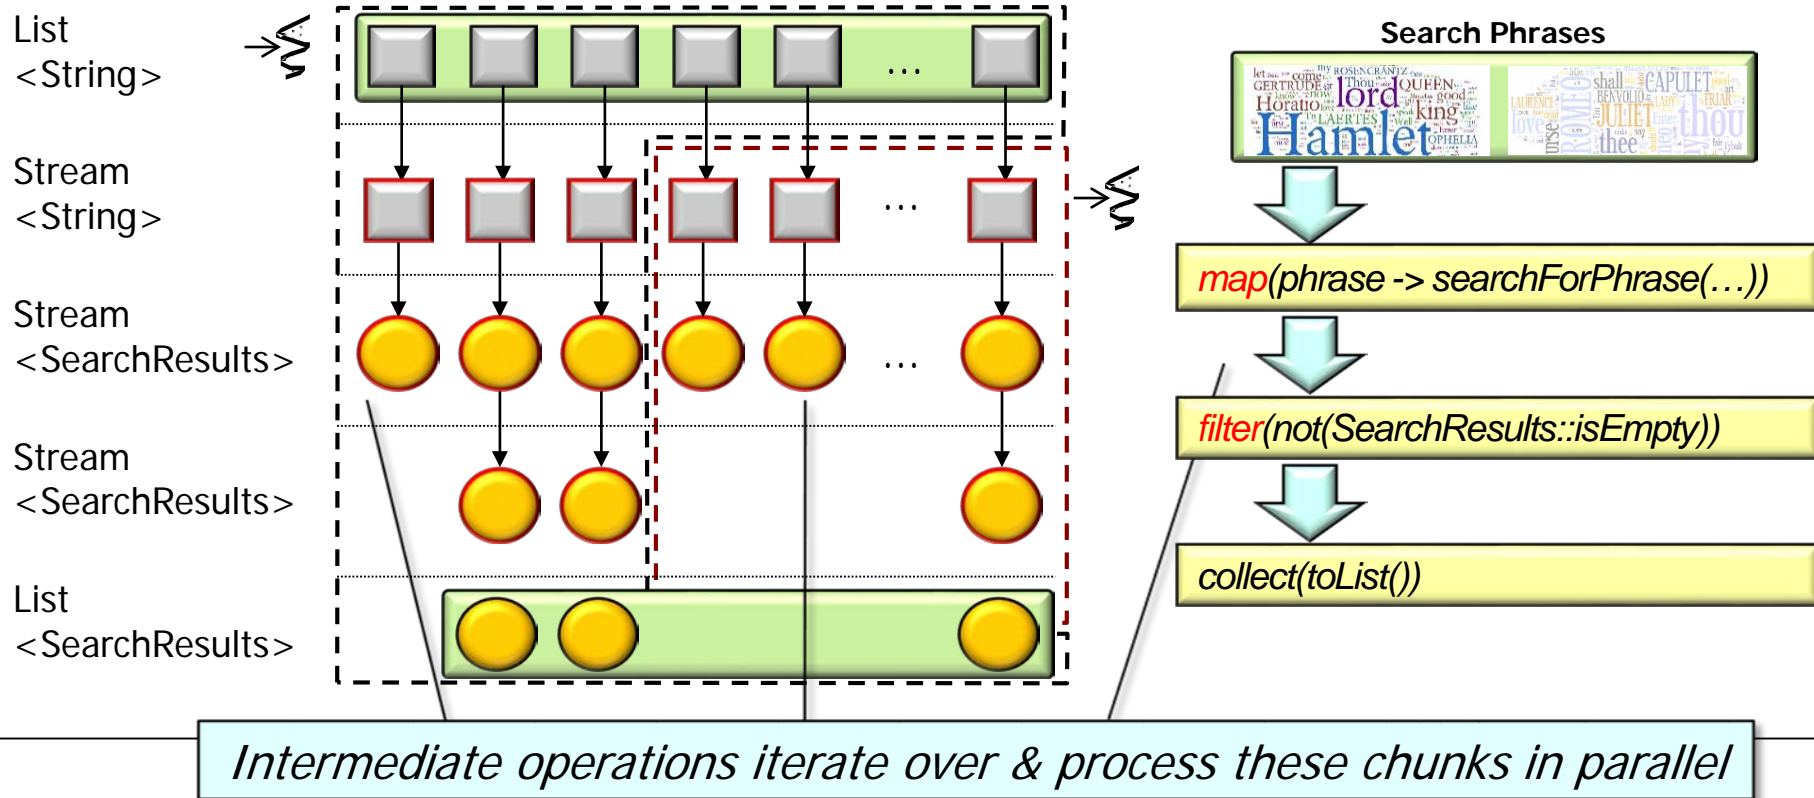
- When a stream executes in parallel, the Java runtime partitions it into multiple substream “chunks” that run in a common fork-join pool



Threads in the pool process different chunks in a non-deterministic order

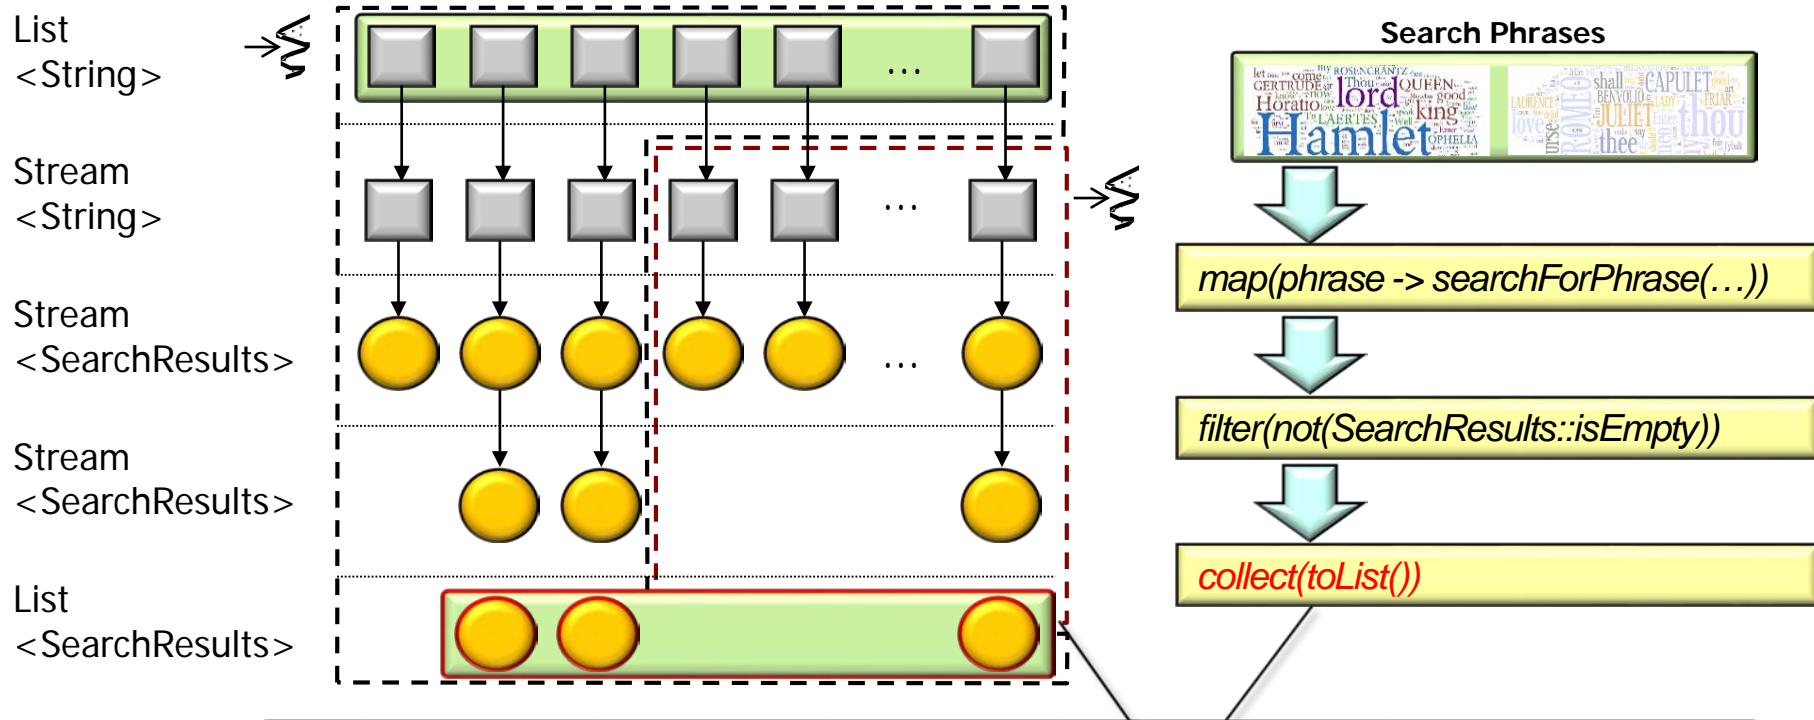
Overview of Java 8 Parallel Streams

- When a stream executes in parallel, the Java runtime partitions it into multiple substream “chunks” that run in a common fork-join pool



Overview of Java 8 Parallel Streams

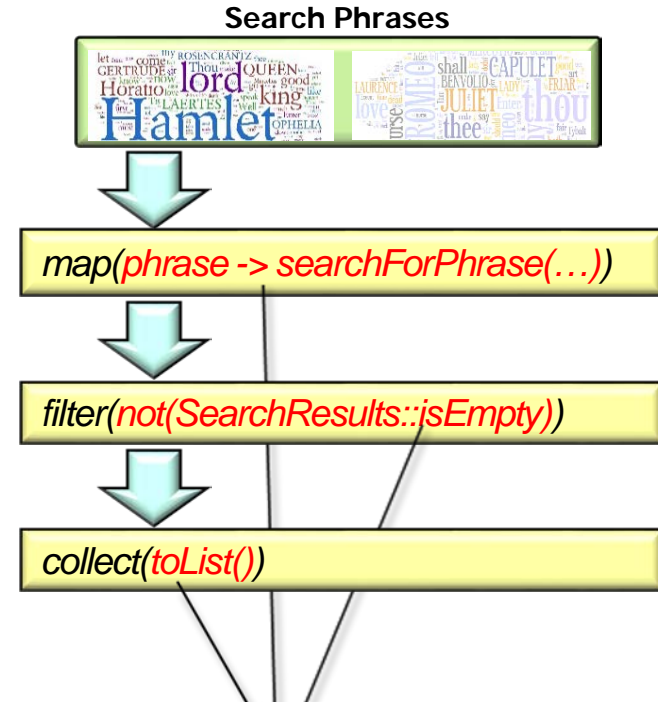

- When a stream executes in parallel, the Java runtime partitions it into multiple substream “chunks” that run in a common fork-join pool



A terminal operation then combines the chunks into a single result

Overview of Java 8 Parallel Streams

- When a stream executes in parallel, the Java runtime partitions it into multiple substream “chunks” that run in a common fork-join pool



(Stateless) Java 8 lambda expressions & method references are used to pass behaviors

Overview of Java 8 Parallel Streams

- The same aggregate operations can be used for sequential & parallel streams

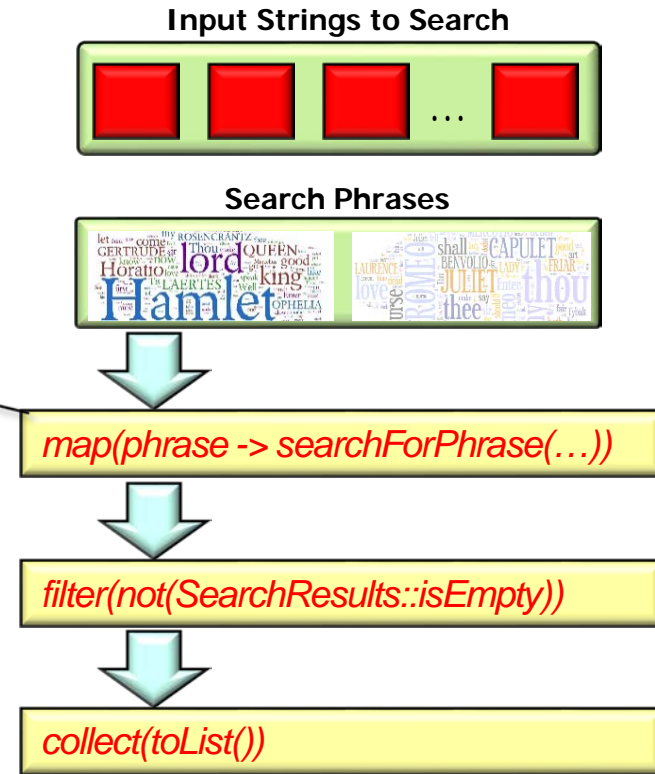
Modifier and Type	Method and Description
boolean	allMatch (Predicate<? super T> predicate) Returns whether all elements of this stream match the provided predicate.
boolean	anyMatch (Predicate<? super T> predicate) Returns whether any elements of this stream match the provided predicate.
static <T> Stream.Builder<T>	builder () Returns a builder for a Stream.
<R,A> R	collect (Collector<? super T,A,R> collector) Performs a mutable reduction operation on the elements of this stream using a Collector.
<R> R	collect (Supplier<R> supplier, BiConsumer<R,? super T> accumulator, BiConsumer<R,R> combiner) Performs a mutable reduction operation on the elements of this stream.
static <T> Stream<T>	concat (Stream<? extends T> a, Stream<? extends T> b) Creates a lazily concatenated stream whose elements are all the elements of the first stream followed by all the elements of the second stream.
long	count () Returns the count of elements in this stream.
Stream<T>	distinct () Returns a stream consisting of the distinct elements (according to <code>Object.equals(Object)</code>) of this stream.
static <T> Stream<T>	empty () Returns an empty sequential Stream.
Stream<T>	filter (Predicate<? super T> predicate) Returns a stream consisting of the elements of this stream that match the given predicate.
Optional<T>	findAny () Returns an Optional describing some element of the stream, or an empty Optional if the stream is empty.
Optional<T>	findFirst () Returns an Optional describing the first element of this stream, or an empty Optional if the stream is empty.
<R> Stream<R>	flatMap (Function<? super T,? extends Stream<? extends R>> mapper) Returns a stream consisting of the results of replacing each element of this stream with the contents of a mapped stream produced by applying the provided mapping function to each element.

See docs.oracle.com/javase/8/docs/api/java/util/stream/Stream.html

Overview of Java 8 Parallel Streams

- The same aggregate operations can be used for sequential & parallel streams



e.g., SearchStreamGang uses the same aggregate operations for both SearchWithSequentialStreams & SearchWithParallelStreams implementations



See github.com/douglasraigschmidt/LiveLessons/tree/master/SearchStreamGang

Overview of Java 8 Parallel Streams

- The same aggregate operations can be used for sequential & parallel streams
- Java 8 streams can thus treat parallelism as an optimization & leverage all available cores!



`map(phrase -> searchForPhrase(...))`



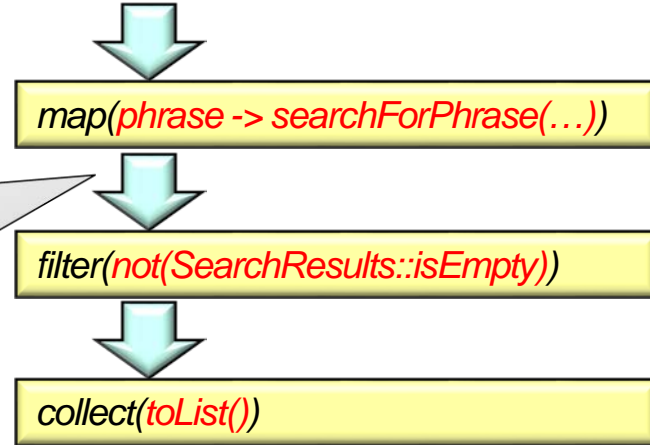

`filter(not(SearchResults::isEmpty))`



`collect(toList())`

Overview of Java 8 Parallel Streams

- The same aggregate operations can be used for sequential & parallel streams
 - Java 8 streams can thus treat parallelism as an optimization & leverage all available cores!
 - Naturally, behaviors run by these aggregate operations must be designed carefully to avoid accessing unsynchronized shared state..

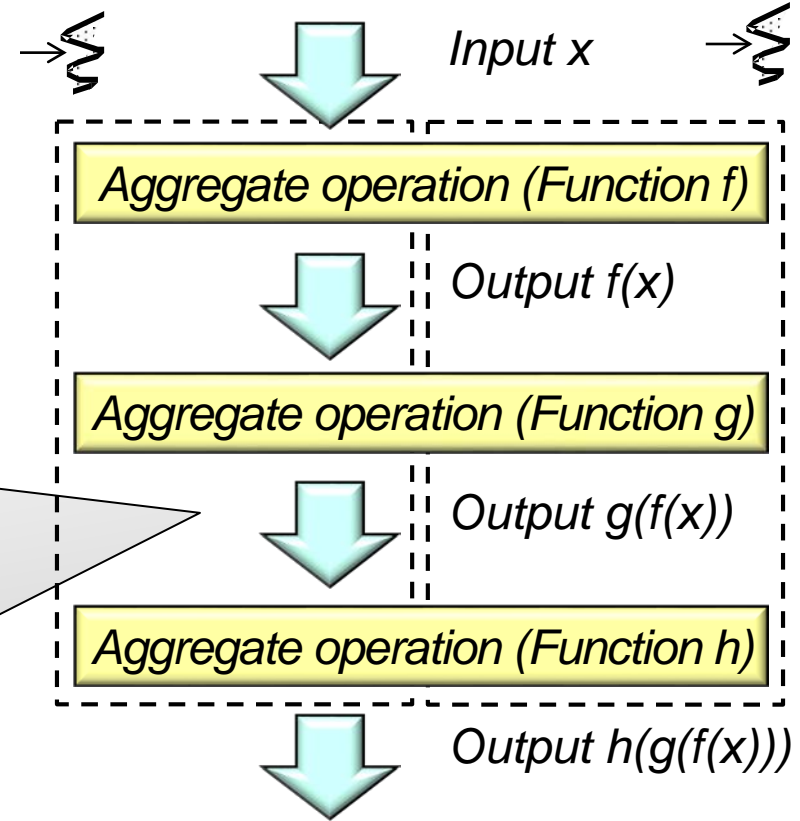


Avoiding Concurrency Hazards in Java 8 Parallel Streams

Avoiding Concurrency Hazards in Java 8 Parallel Streams

- Java 8 parallel streams assume behaviors incur no race conditions

Race conditions arise when an app depends on the sequence or timing of threads for it to operate properly



Avoiding Concurrency Hazards in Java 8 Parallel Streams


- Parallel streams should therefore avoid operations with side-effects



Input Strings to Search



Search Phrases



```
map(phrase -> searchForPhrase(...))
```



```
filter(not(SearchResults::isEmpty))
```



```
collect(toList())
```

See docs.oracle.com/javase/tutorial/collections/streams/parallelism.html#side_effects

Avoiding Concurrency Hazards in Java 8 Parallel Streams

- Parallel streams should therefore avoid operations with side-effects, e.g.
- *Stateful lambda expressions*
 - Where results depends on state that may change in concurrent execution of a pipeline

```
class BuggyFactorial {  
    static class Total {  
        long mTotal = 1;  
        void multiply(long n)  
        { mTotal *= n; }  
    }  
  
    static long factorial(long n){  
        Total t = new Total();  
        LongStream  
            .rangeClosed(1, n)  
            .parallel()  
            .forEach(t::multiply);  
  
        return t.mTotal;  
    } ...  
}
```

Avoiding Concurrency Hazards in Java 8 Parallel Streams

- Parallel streams should therefore avoid operations with side-effects, e.g.
- Stateful lambda expressions*
 - Where results depends on state that may change in concurrent execution of a pipeline

Race conditions can arise due to the unsynchronized access to mTotal field

```
class BuggyFactorial {  
    static class Total {  
        long mTotal = 1;  
        void multiply(long n)  
        { mTotal *= n; }  
    }  
  
    static long factorial(long n){  
        Total t = new Total();  
        LongStream  
            .rangeClosed(1, n)  
            .parallel()  
            .forEach(t::multiply);  
  
        return t.mTotal;  
    } ...  
}
```

Avoiding Concurrency Hazards in Java 8 Parallel Streams

- Parallel streams should therefore avoid operations with side-effects, e.g.

- *Stateful lambda expressions*

- *Interference w/the data source*

- Occurs when the source of a stream is modified while a pipeline processes the stream

```
List<Integer> list = IntStream  
    .range(0, 10)  
    .boxed()  
    .collect(toList());
```

```
list  
    .parallelStream()  
    .peek(list::remove)  
    .forEach(System.out::println);
```


Avoiding Concurrency Hazards in Java 8 Parallel Streams

- Parallel streams should therefore avoid operations with side-effects, e.g.

- Stateful lambda expressions*
- Interference w/the data source*
 - Occurs when the source of a stream is modified while a pipeline processes the stream

```
List<Integer> list = IntStream  
    .range(0, 10)  
    .boxed()  
    .collect(toList());
```

```
list  
    .parallelStream()  
    .peek(list::remove)  
    .forEach(System.out::println);
```

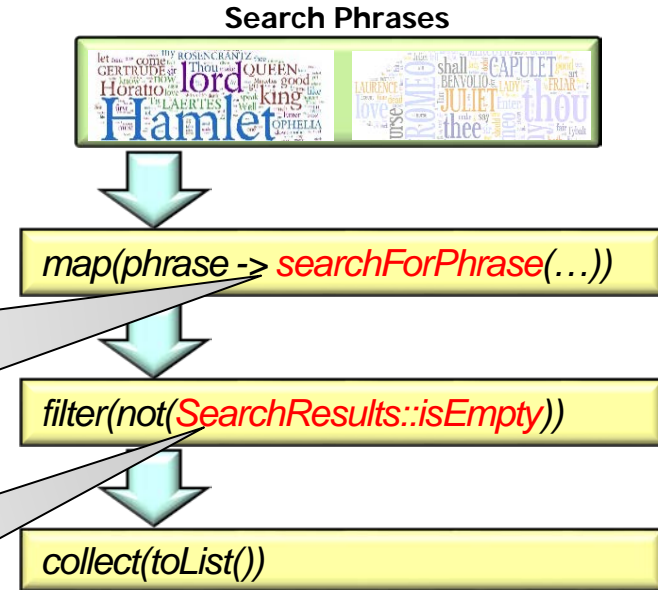
Aggregate operations enable parallelism with non-thread-safe collections provided the collection is not modified while it's being operated on..

Avoiding Concurrency Hazards in Java 8 Parallel Streams

- Java 8 lambda expressions & method references containing no shared state are useful for parallel streams since they needn't be explicitly synchronized

```
return new SearchResults  
(Thread.currentThread().getId(),  
currentCycle(), phrase, title,  
StreamSupport  
    .stream(new PhraseMatchSpliterator  
        (input, phrase),  
        parallel)  
    .collect(toList()));
```

```
return mList.size() == 0;
```



End of Overview of Java 8 Parallel Streams (Part 1)