

Overview of Java 8 Parallel Streams

(Part 2)

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

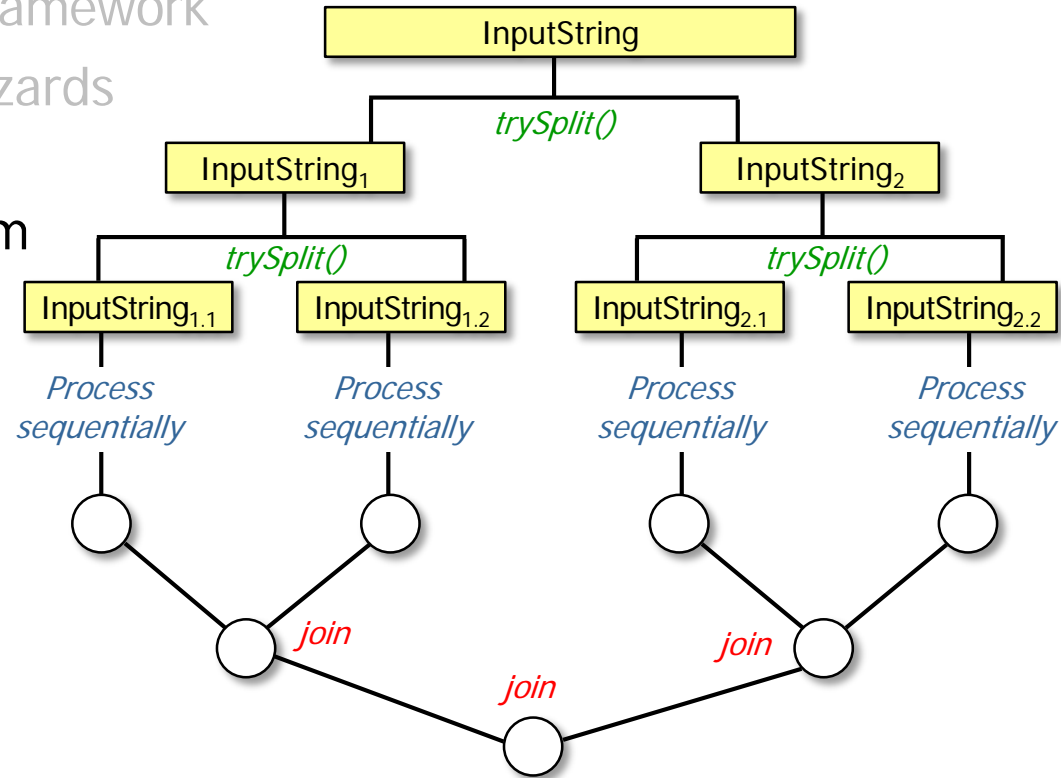
Institute for Software
Integrated Systems

Vanderbilt University
Nashville, Tennessee, USA



Learning Objectives in this Part of the Lesson

- Recognize how Java 8 applies aggregate operations & functional programming features in the parallel streams framework
- Be able to avoid concurrency hazards in parallel streams
- Understand how a parallel stream splits its elements recursively, processes them independently, & combines the results




This knowledge of internals will make you a better Java 8 streams programmer!

Inner Workings of a Parallel Stream

Inner Workings of a Parallel Stream

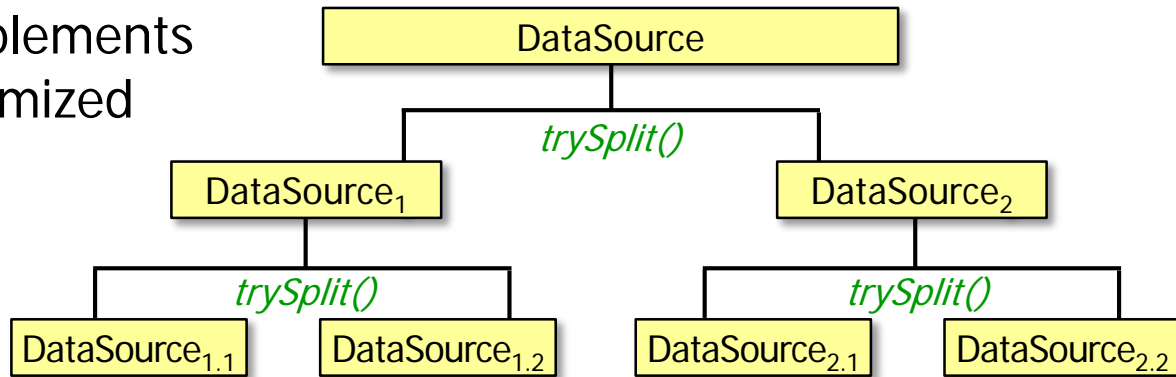
- A Java 8 parallel stream implements a “map/reduce” variant optimized for multi-core processors



See en.wikipedia.org/wiki/MapReduce

Inner Workings of a Parallel Stream

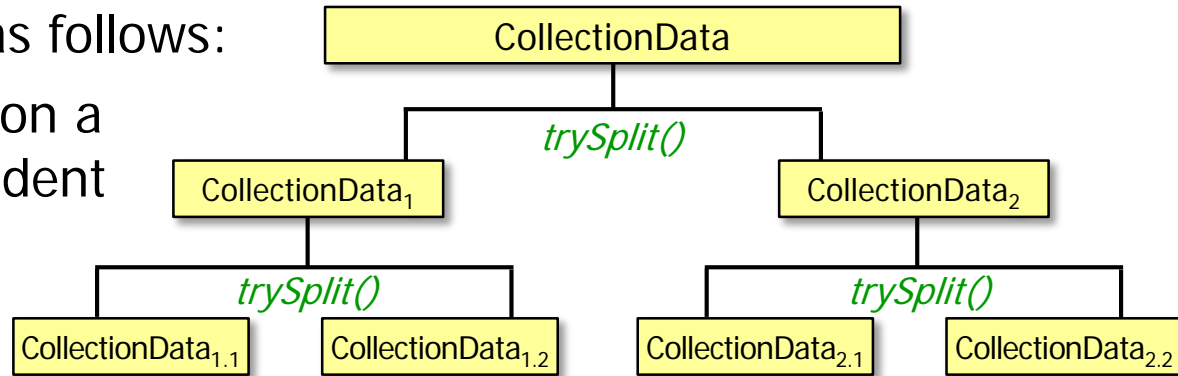
- A Java 8 parallel stream implements a “map/reduce” variant optimized for multi-core processors
- It’s actually more like the “split-apply-combine” data analysis strategy



Inner Workings of a Parallel Stream

- Split-apply-combine works as follows:

1. Split – Recursively partition a data source into independent “chunks”



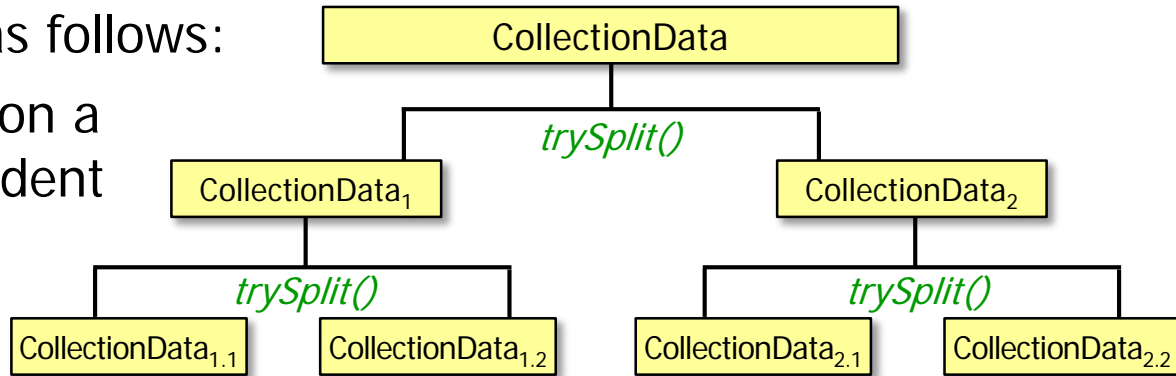
See en.wikipedia.org/wiki/Divide_and_conquer_algorithm

Inner Workings of a Parallel Stream

- Split-apply-combine works as follows:

1. Split – Recursively partition a data source into independent “chunks”

- Splitterators are defined to partition collections in Java 8



```
public interface Splitterator<T> {  
    boolean tryAdvance(Consumer<? super T> action) ;  
    Splitterator<T> trySplit() ;  
    long estimateSize();  
    int characteristics();  
}
```

See docs.oracle.com/javase/8/docs/api/java/util/Splitterator.html


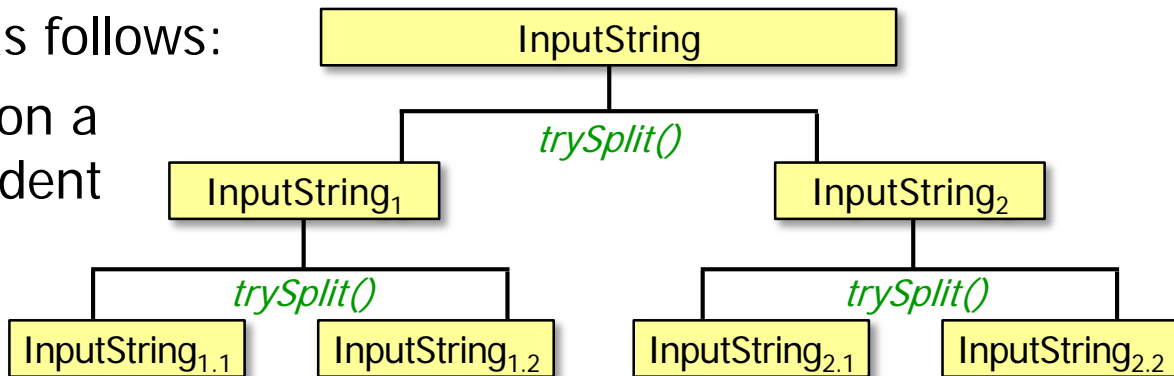
Inner Workings of a Parallel Stream

- Split-apply-combine works as follows:

1. Split – Recursively partition a data source into independent “chunks”

- Spliterators are defined to partition collections in Java 8

- You can also define custom spliterators

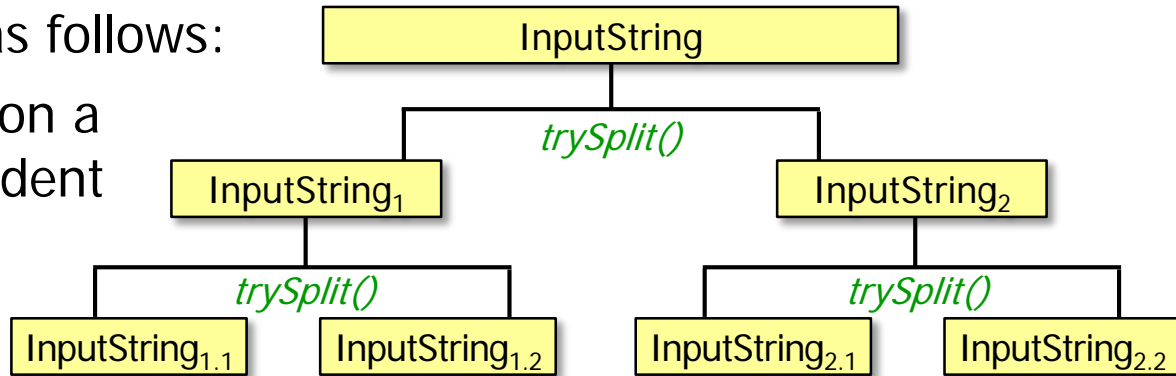


Inner Workings of a Parallel Stream

- Split-apply-combine works as follows:

1. Split – Recursively partition a data source into independent “chunks”

- Splitterators are defined to partition collections in Java 8
- You can also define custom splitterators
- Parallel streams perform better on data sources that can be split efficiently & evenly

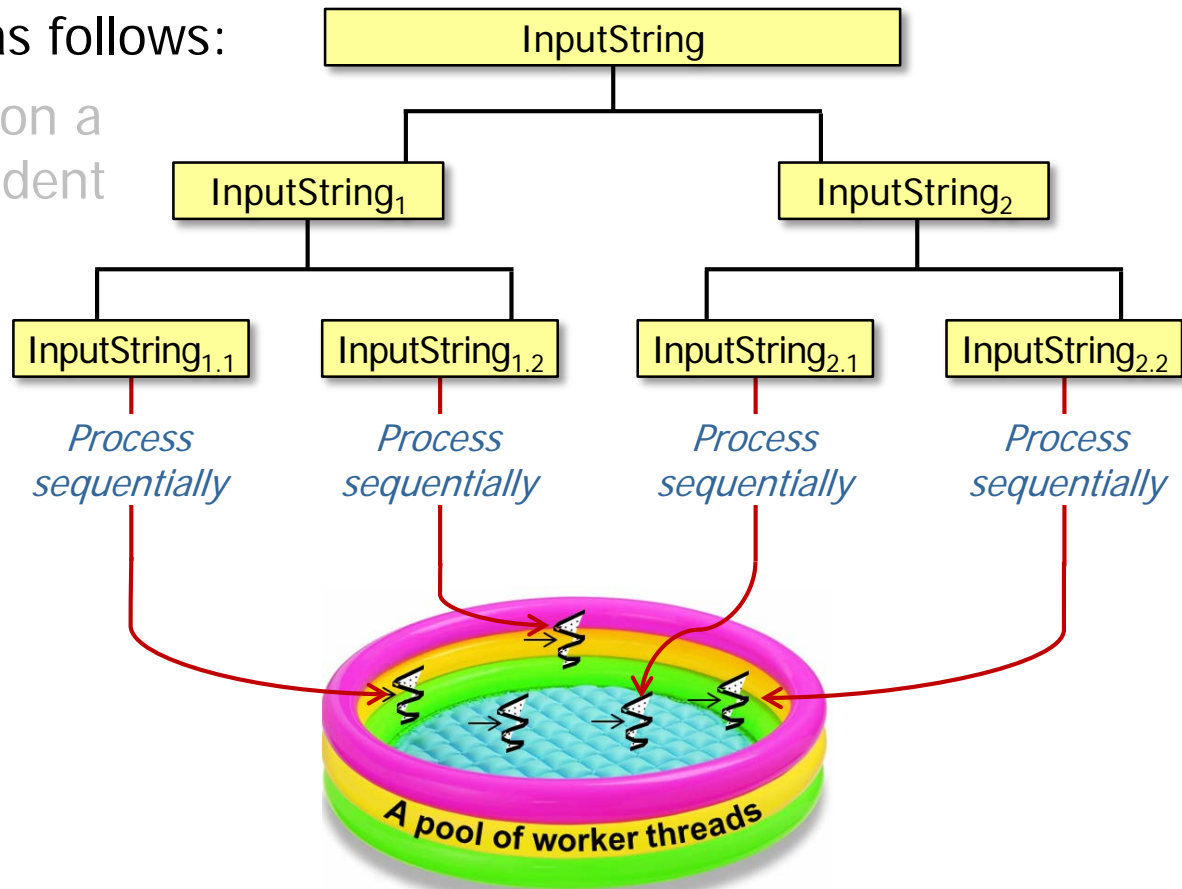


Inner Workings of a Parallel Stream

- Split-apply-combine works as follows:

1. Split – Recursively partition a data source into independent “chunks”

2. Apply – Process chunks independently in a pool of threads

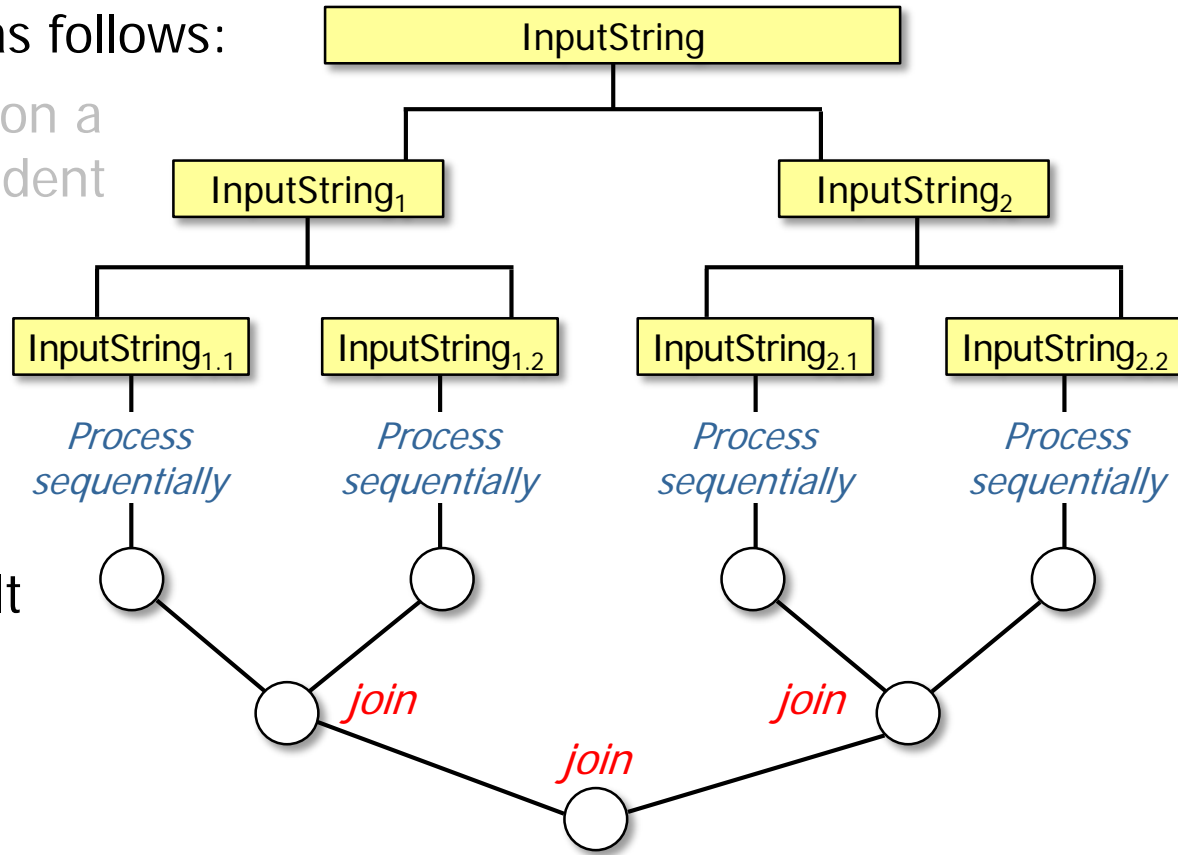


In practice, splitting & applying *run* simultaneously, not sequentially

Inner Workings of a Parallel Stream

- Split-apply-combine works as follows:

1. Split – Recursively partition a data source into independent “chunks”
2. Apply – Process chunks independently in a pool of threads
3. Combine – Join partial results into a single result



Combining is performed by terminal operations, such as `collect()` & `reduce()`

Partitioning a Parallel Stream

Partitioning a Parallel Stream

- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks

Interface `Spliterator<T>`

Type Parameters:

`T` - the type of elements returned by this `Spliterator`

All Known Subinterfaces:

`Spliterator.OfDouble`, `Spliterator.OfInt`, `Spliterator.OfLong`,
`Spliterator.OfPrimitive<T,T_CONS,T_SPLITR>`

All Known Implementing Classes:

`Spliterators.AbstractDoubleSpliterator`,
`Spliterators.AbstractIntSpliterator`,
`Spliterators.AbstractLongSpliterator`,
`Spliterators.AbstractSpliterator`

```
public interface Spliterator<T>
```

An object for traversing and partitioning elements of a source. The source of elements covered by a `Spliterator` could be, for example, an array, a `Collection`, an IO channel, or a generator function.

A `Spliterator` may traverse elements individually (`tryAdvance()`) or sequentially in bulk (`forEachRemaining()`).

See docs.oracle.com/javase/8/docs/api/java/util/Spliterator.html

Partitioning a Parallel Stream

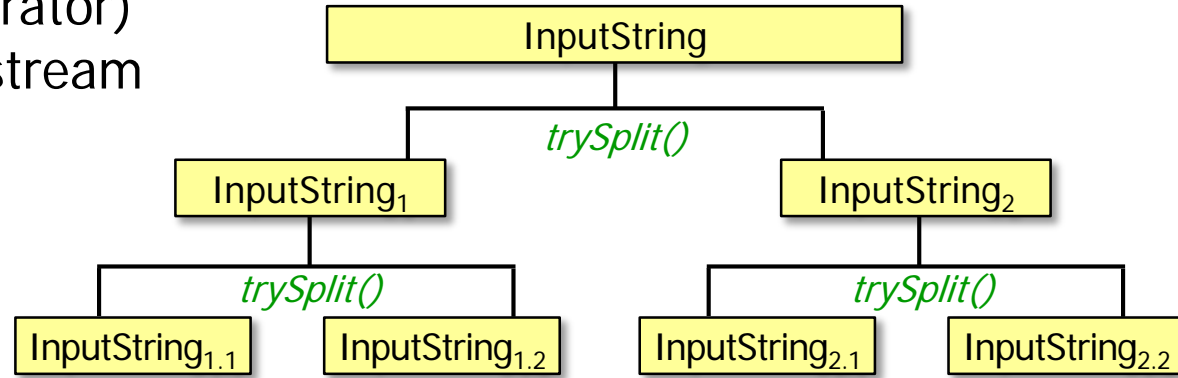
- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks
- We showed earlier how a spliterator can *traverse* elements in a source

```
List<String> quote = Arrays.asList  
    ("This ", "above ", "all- ",  
     "to ", "thine ", "own ",  
     "self ", "be ", "true", ",\n",  
     ...);  
  
for(Spliterator<String> s =  
    quote.spliterator();  
    s.tryAdvance(System.out::print)  
        != false;  
    )  
    continue;
```

See “*Overview of Java Streams (Part 3)*”

Partitioning a Parallel Stream

- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks
 - We showed earlier how a spliterator can *traverse* elements in a source
- We now outline how a parallel spliterator can *partition* all elements in a source

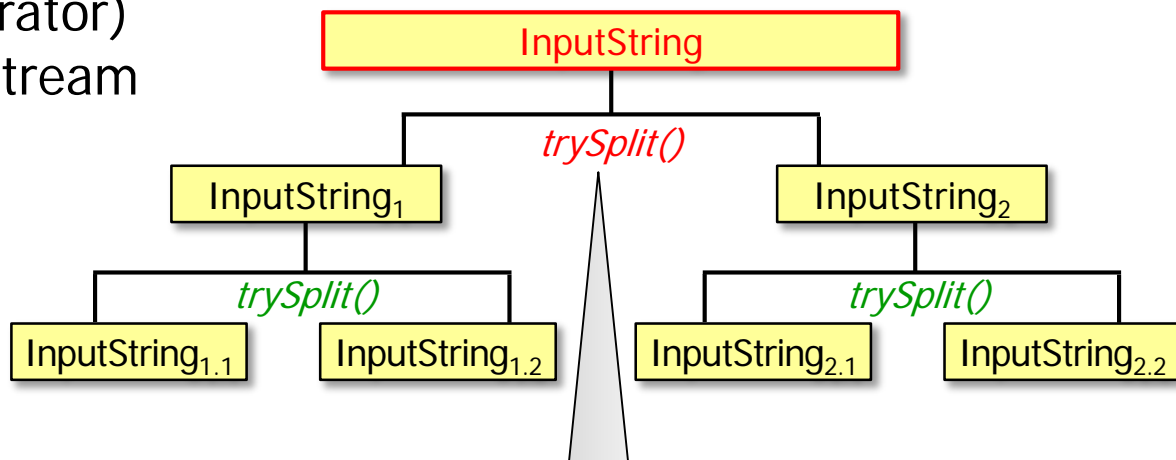


Partitioning a Parallel Stream

- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks

- We showed earlier how a spliterator can *traverse* elements in a source

- We now outline how a parallel spliterator can *partition* all elements in a source

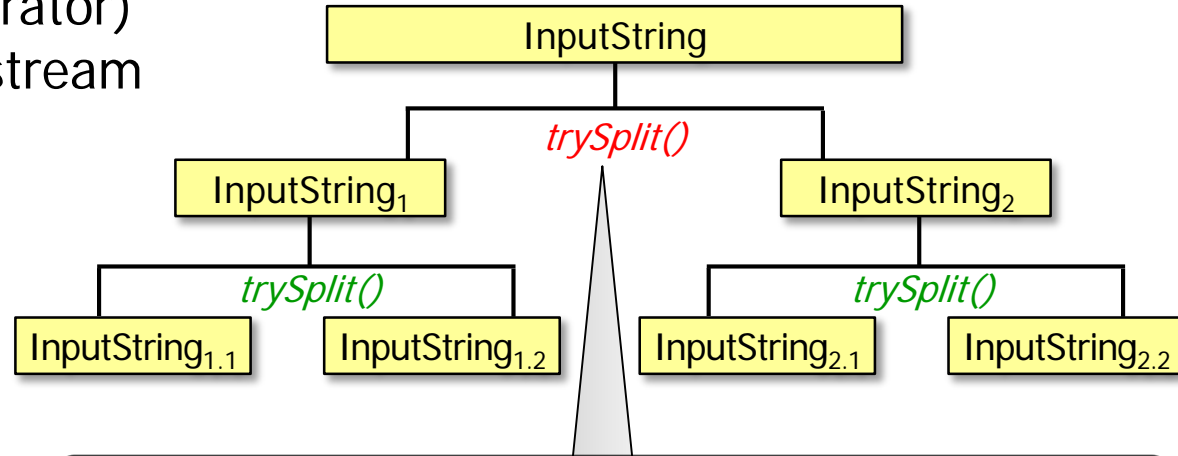


```
Spliterator<T> trySplit() {  
    if (input is below minimum size)  
        return null  
    else {  
        split input in 2 (even-sized) chunks  
        return a spliterator for "left chunk"  
    }  
}
```

This partitioning is done via a spliterator's trySplit() method

Partitioning a Parallel Stream

- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks
 - We showed earlier how a spliterator can *traverse* elements in a source
 - We now outline how a parallel spliterator can *partition* all elements in a source

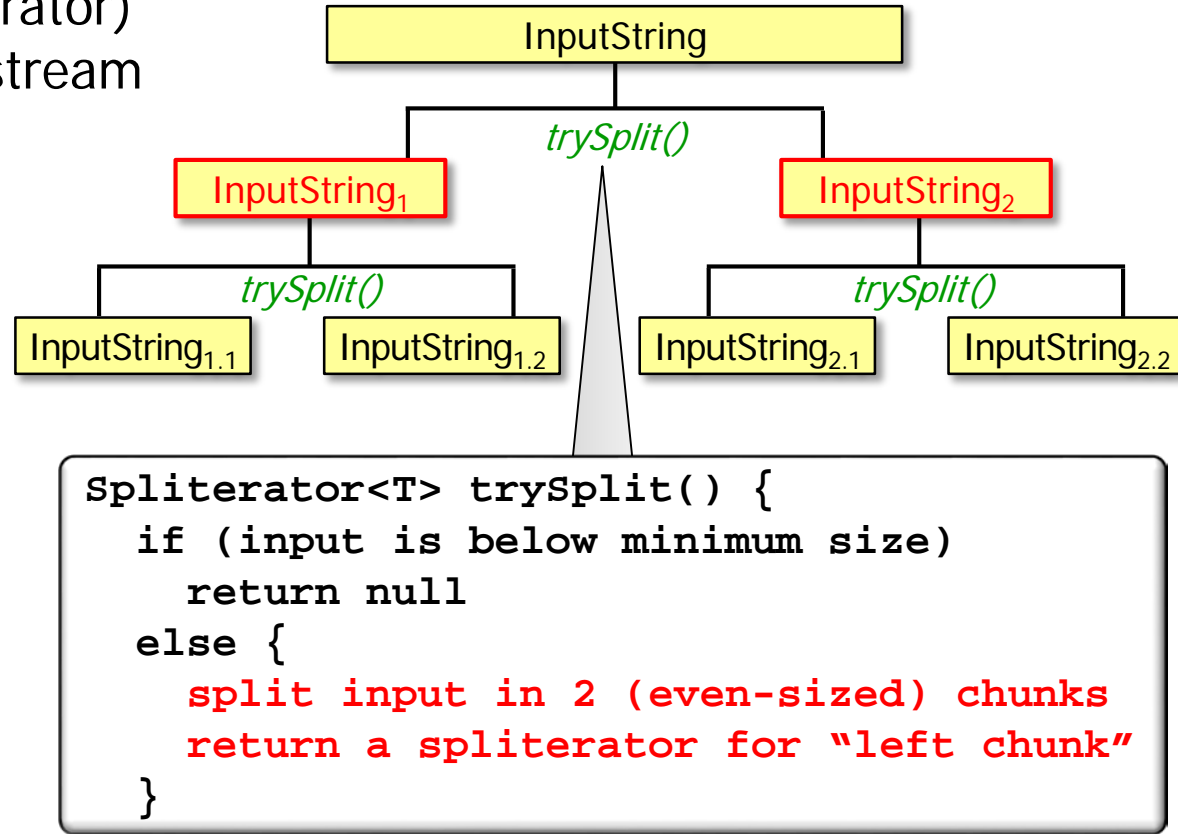


```
Spliterator<T> trySplit() {  
    if (input is below minimum size)  
        return null  
    else {  
        split input in 2 (even-sized) chunks  
        return a spliterator for "left chunk"  
    }  
}
```

When null is returned the streams framework processes this chunk sequentially

Partitioning a Parallel Stream

- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks
 - We showed earlier how a spliterator can *traverse* elements in a source
 - We now outline how a parallel spliterator can *partition* all elements in a source



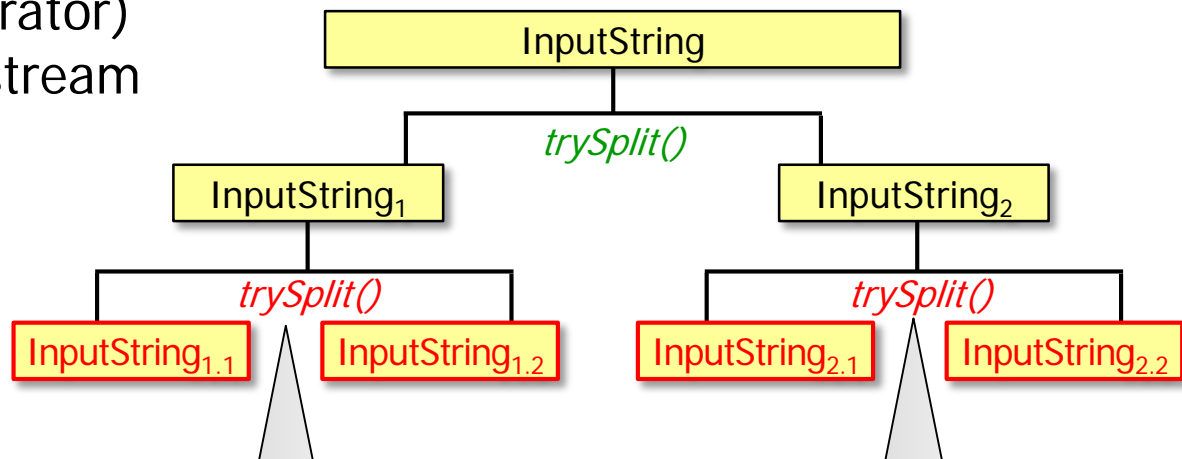
A spliterator usually needs no synchronization nor does it need a “join” phase!

Partitioning a Parallel Stream

- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks

- We showed earlier how a spliterator can *traverse* elements in a source

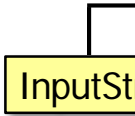
- We now outline how a parallel spliterator can *partition* all elements in a source


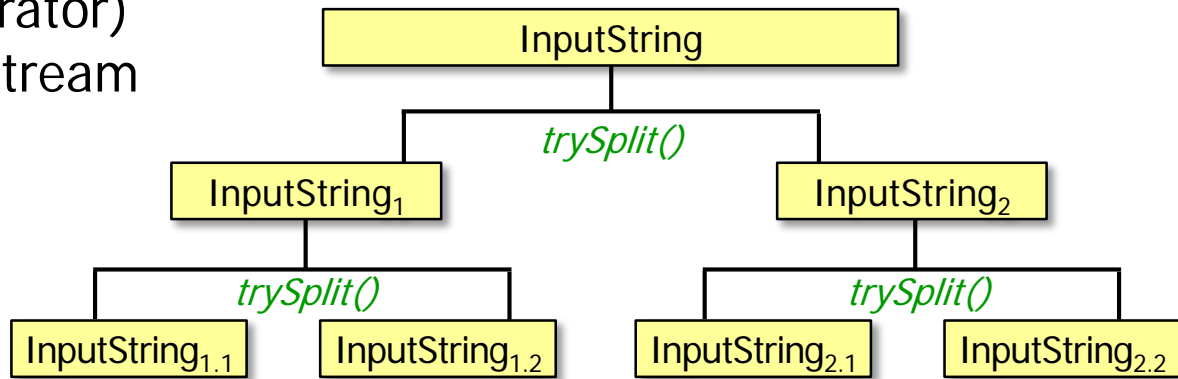


```
Spliterator<T> trySplit() {  
    if (input is below minimum size)  
        return null  
    else {  
        split input in 2 (even-sized) chunks  
        return a spliterator for "left chunk"  
    }  
}
```

trySplit() is called recursively until all chunks are below the minimize size

Partitioning a Parallel Stream

- A “splittable iterator” (spliterator) partitions a Java 8 parallel stream into chunks
 - We showed earlier how a spliterator can *traverse* elements in a source
 - We now outline how a parallel spliterator can *partition* all elements in a source
 - More parallel spliterator implementation details are covered shortly
- 




See *"Java 8 Parallel SearchStreamGang Example (Part 2)"*

Processing Chunks in a Parallel Stream


Processing Chunks in a Parallel Stream

- The chunks created by a spliterator are processed in a common fork-join pool



Processing Chunks in a Parallel Stream

- The chunks created by a spliterator are processed in a common fork-join pool
- All parallel streams in a process share this common fork-join pool



Processing Chunks in a Parallel Stream

- ForkJoinPool is an Executor Service implementation that runs ForkJoinTasks

Class ForkJoinPool

```
java.lang.Object  
    java.util.concurrent.AbstractExecutorService  
        java.util.concurrent.ForkJoinPool
```

All Implemented Interfaces:

```
Executor, ExecutorService
```

```
public class ForkJoinPool  
    extends AbstractExecutorService
```

An `ExecutorService` for running `ForkJoinTasks`. A `ForkJoinPool` provides the entry point for submissions from non-`ForkJoinTask` clients, as well as management and monitoring operations.

A `ForkJoinPool` differs from other kinds of `ExecutorService` mainly by virtue of employing *work-stealing*: all threads in the pool attempt to find and execute tasks submitted to the pool and/or created by other active tasks (eventually blocking waiting for work if none exist). This enables efficient processing when most tasks spawn other subtasks (as do most `ForkJoinTasks`), as well as when many small tasks are submitted to the pool from external clients. Especially when setting *asyncMode* to true in constructors, `ForkJoinPools` may also be appropriate for use with event-style tasks that are never joined.

A static `commonPool()` is available and appropriate for most applications. The common pool is used by any `ForkJoinTask` that is not explicitly submitted to a specified pool. Using the common pool normally reduces resource usage (its threads are slowly reclaimed during periods of non-use, and reinstated upon subsequent use).

See docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinPool.html

Processing Chunks in a Parallel Stream

- ForkJoinPool is an Executor Service implementation that runs ForkJoinTasks
 - It provides the entry point for submissions from non-ForkJoinTask clients

void	<u>execute(ForkJoinTask<T>)</u> – Arrange async execution
T	<u>invoke(ForkJoinTask<T>)</u> – Performs the given task, returning its result upon completion
<u>ForkJoinTask<T></u>	<u>submit(ForkJoinTask)</u> – Submits a ForkJoinTask for execution, returns a future

See docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinPool.html

Processing Chunks in a Parallel Stream

- ForkJoinPool is an Executor Service implementation that runs ForkJoinTasks
 - It provides the entry point for submissions from non-ForkJoinTask clients
 - It also provides management & monitoring operations

int	getParallelism() – Returns the targeted parallelism level of this pool
int	getPoolSize() – Returns the number of worker threads that have started but not yet terminated
int	getQueuedSubmissionCount() – Returns an estimate of the number of tasks submitted to this pool that have not yet begun executing
long	getStealCount() – Returns an estimate of the total number of tasks stolen from one thread's work queue by another

See docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinPool.html

Processing Chunks in a Parallel Stream

- A ForkJoinTask is a chunk of data along with functionality on that data

Class ForkJoinTask<V>

java.lang.Object
java.util.concurrent.ForkJoinTask<V>

All Implemented Interfaces:

Serializable, Future<V>

Direct Known Subclasses:

CountedCompleter, RecursiveAction, RecursiveTask

```
public abstract class ForkJoinTask<V>  
extends Object  
implements Future<V>, Serializable
```

Abstract base class for tasks that run within a `ForkJoinPool`. A `ForkJoinTask` is a thread-like entity that is much lighter weight than a normal thread. Huge numbers of tasks and subtasks may be hosted by a small number of actual threads in a `ForkJoinPool`, at the price of some usage limitations.

A "main" `ForkJoinTask` begins execution when it is explicitly submitted to a `ForkJoinPool`, or, if not already engaged in a `ForkJoin` computation, commenced in the `ForkJoinPool.commonPool()` via `fork()`, `invoke()`, or related methods. Once started, it will usually in turn start other subtasks. As indicated by the name of this class, many programs using `ForkJoinTask` employ only methods `fork()` and `join()`, or derivatives such as `invokeAll`. However, this class also provides a number of other methods that can come into play in advanced usages, as well as extension mechanics that allow support of new forms of fork/join processing.

See docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinTask.html

Processing Chunks in a Parallel Stream

- A ForkJoinTask is a chunk of data along with functionality on that data
 - It defines two primary methods

[ForkJoinTask](#)

<T>


V

[fork\(\)](#) – Arranges to asynchronously execute this task in the appropriate pool


[join\(\)](#) – Returns the result of the computation when it is done

Processing Chunks in a Parallel Stream

- A ForkJoinTask is a chunk of data along with functionality on that data
 - It defines two primary methods
- A ForkJoinTask is lighter weight than a Java thread




ForkJoinTask



Thread

Processing Chunks in a Parallel Stream


- A ForkJoinTask is a chunk of data along with functionality on that data
 - It defines two primary methods
 - A ForkJoinTask is lighter weight than a Java thread
 - A large # of ForkJoinTasks can run in a small # of Java threads in a ForkJoinPool



See www.infoq.com/interviews/doug-lea-fork-join

Processing Chunks in a Parallel Stream


- A circular deque is associated with each ForkJoinPool thread



See en.wikipedia.org/wiki/Double-ended_queue


Processing Chunks in a Parallel Stream

- A circular deque is associated with each ForkJoinPool thread
- `fork()` pushes a new task to the head of its deque




Processing Chunks in a Parallel Stream

- A circular deque is associated with each ForkJoinPool thread
 - `fork()` pushes a new task to the head of its deque
- Likewise, a thread pops the next task its processes from the head of its deque




Processing Chunks in a Parallel Stream

- A circular deque is associated with each ForkJoinPool thread
 - `fork()` pushes a new task to the head of its deque
 - Likewise, a thread pops the next task its processes from the head of its deque
- An idle thread “steals” work from the tail of a busy thread to maximize core utilization




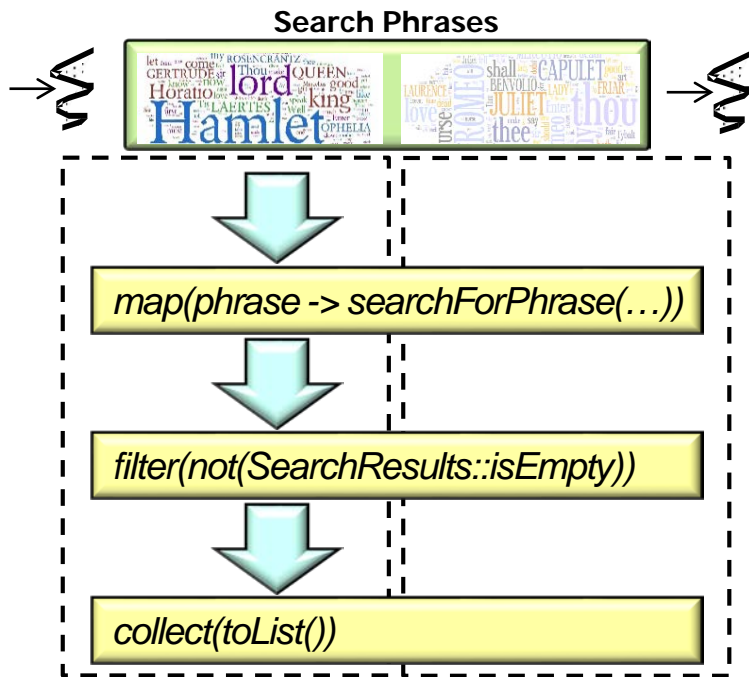
Processing Chunks in a Parallel Stream

- A circular deque is associated with each ForkJoinPool thread
 - `fork()` pushes a new task to the head of its deque
 - Likewise, a thread pops the next task its processes from the head of its deque
- An idle thread “steals” work from the tail of a busy thread to maximize core utilization
 - The circular deque used for work-stealing lowers contention



Processing Chunks in a Parallel Stream

- Parallel streams is a “user friendly” ForkJoinPool façade



See espressoprogrammer.com/fork-join-vs-parallel-stream-java-8

Processing Chunks in a Parallel Stream

- Parallel streams is a “user friendly” ForkJoinPool façade
- You can program directly to the ForkJoinPool API, though it can be somewhat tedious!

```
List<List<SearchResults>>  
listOfListOfSearchResults =  
    ForkJoinPool  
        .commonPool()  
        .invoke(new  
            SearchWithForkJoinTask  
                (inputList,  
                 mPhrasesToFind,  
                 ...));
```

Processing Chunks in a Parallel Stream

- Parallel streams is a “user friendly” ForkJoinPool façade
- You can program directly to the ForkJoinPool API, though it can be somewhat tedious!

Use the common fork-join pool to search input strings for phrases that match

```
List<List<SearchResults>>  
    listOfListOfSearchResults =  
        ForkJoinPool  
            .commonPool()  
            .invoke(new  
                SearchWithForkJoinTask  
                    (inputList,  
                     mPhrasesToFind,  
                     ...));
```

Input Strings to Search



Search Phrases



Processing Chunks in a Parallel Stream

- Parallel streams is a “user friendly” ForkJoinPool façade
- You can program directly to the ForkJoinPool API, though it can be somewhat tedious!
- Mostly used for parallel algorithms that don't match Java 8's parallel streams programming model

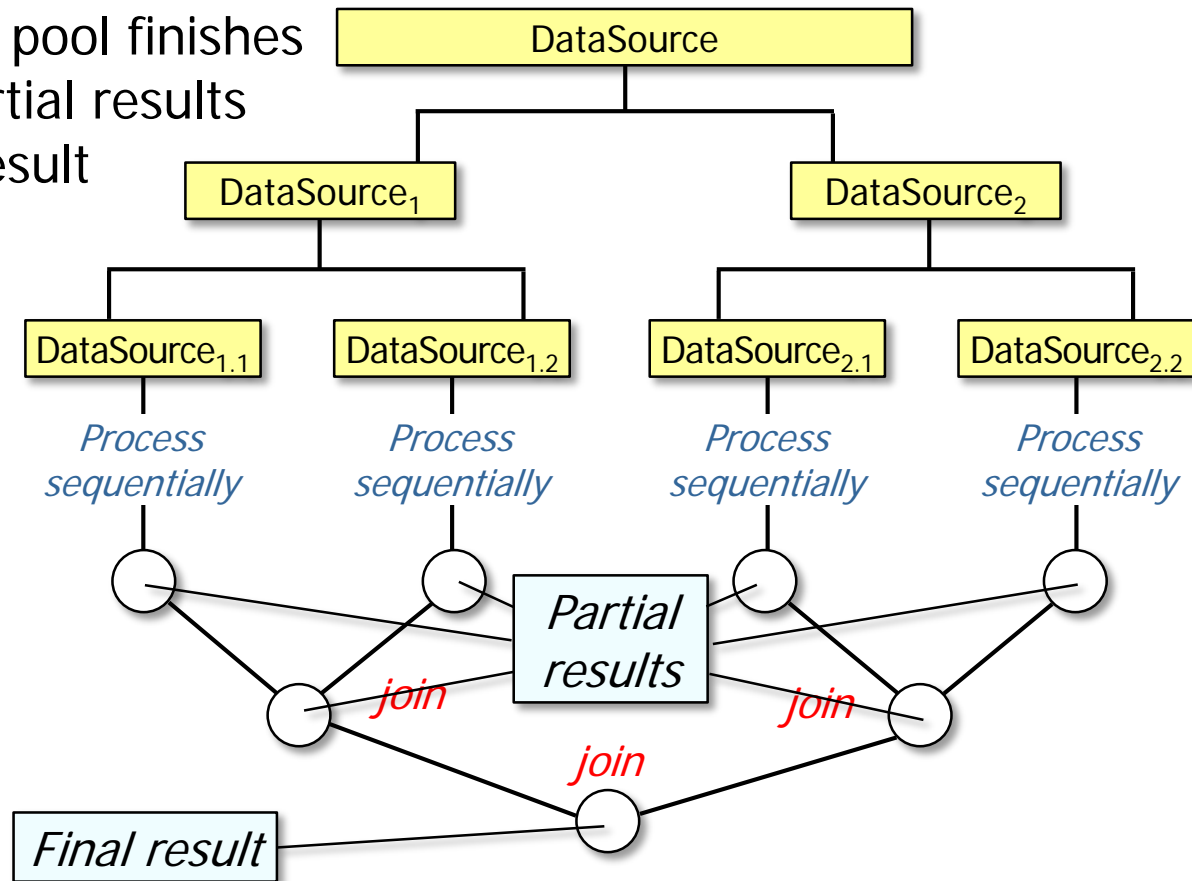
```
if (n % 2 == 1) {  
    //  $f(2n-1) = f(n-1)^2 + f(n)^2$   
    int left = (n + 1) / 2;  
    int right = (n + 1) / 2 - 1;  
    FibonacciTask f0 =  
        new FibonacciTask(left);  
    FibonacciTask f1 =  
        new FibonacciTask(right);  
    f1.fork();  
    BigInteger bi0 = f0.invoke();  
    BigInteger bi1 = f1.join();  
    if (isCancelled()) return null;  
    result =  
        square(bi1).add(square(bi0));  
} else {  
    //  $f(2n) = (2*f(n-1)+f(n))*f(n)$ 
```

See www.javaspecialists.eu/archive/Issue201.html

Combining Results in a Parallel Stream

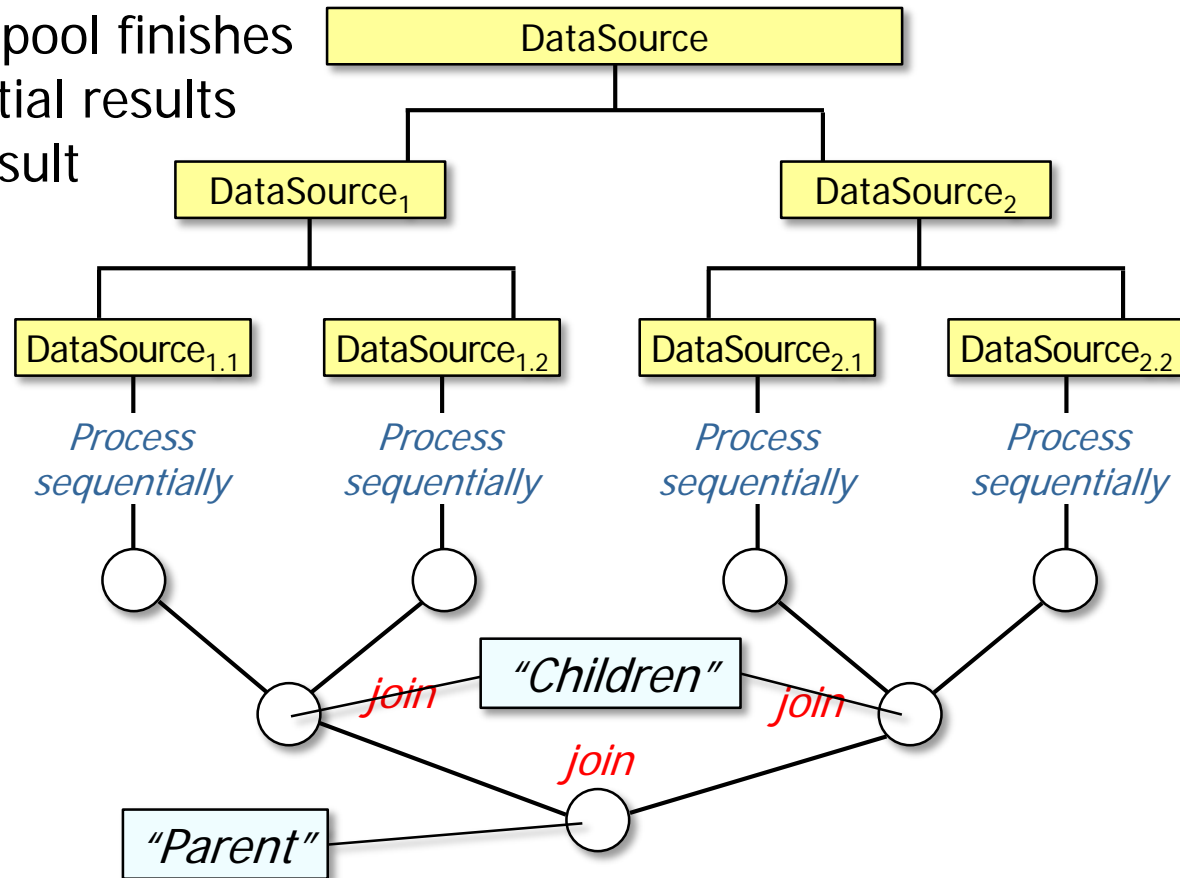
Combining Results in a Parallel Stream

- After the common fork-join pool finishes processing chunks their partial results are combined into a final result



Combining Results in a Parallel Stream


- After the common fork-join pool finishes processing chunks their partial results are combined into a final result
 - join() occurs in a single thread at each level
 - i.e., the “parent”



As a result, there's typically no need for synchronizers during the joining

Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways



Understanding these differences is particularly important for parallel streams

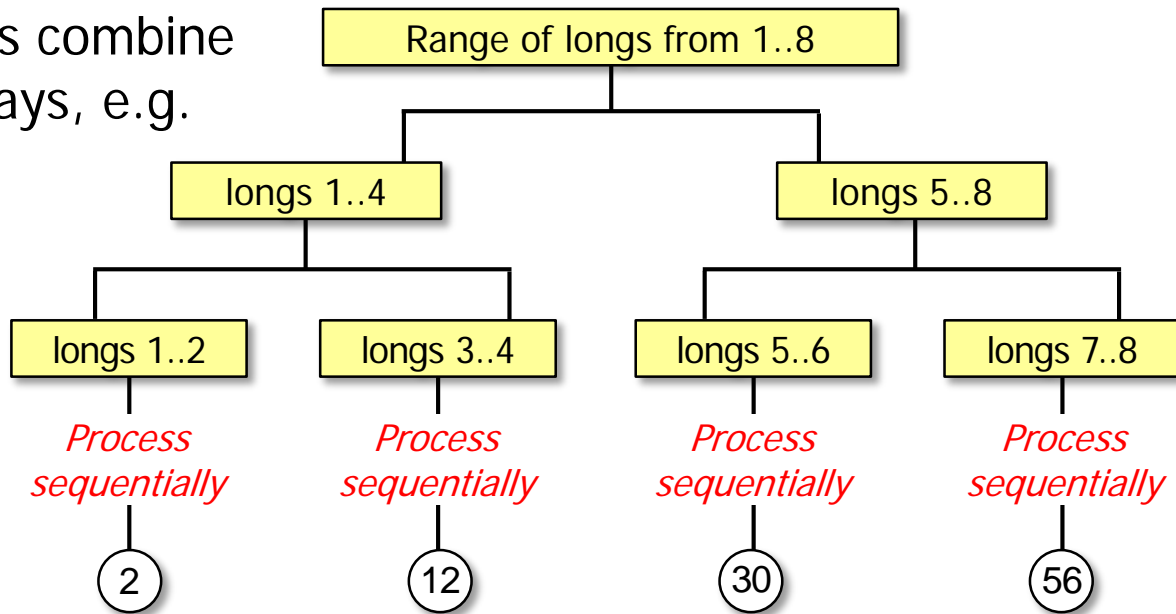
Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
 - `reduce()` creates a new immutable value



Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
- `reduce()` creates a new immutable value

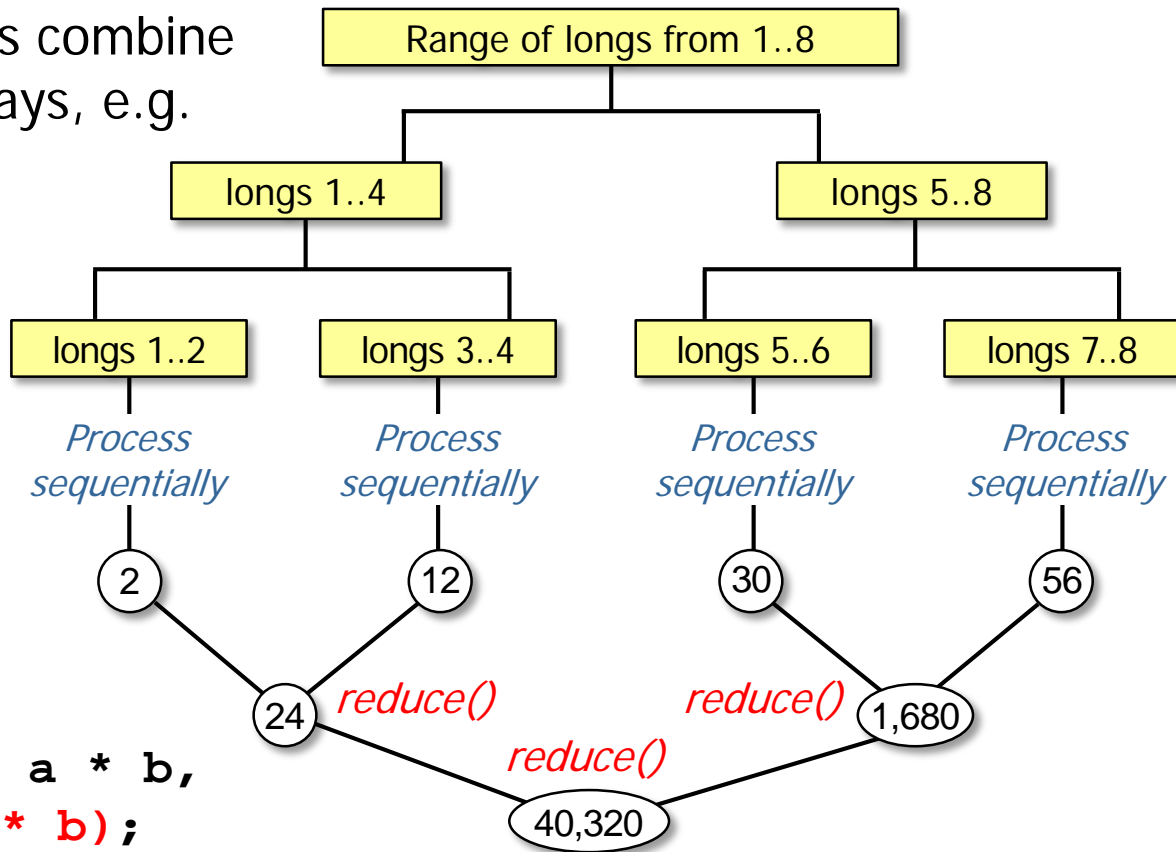


```
long factorial(long n) {  
    return LongStream  
        .rangeClosed(1, n)  
        .parallel()  
        .reduce(1, (a, b) -> a * b,  
                (a, b) -> a * b);  
}
```

Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
- `reduce()` creates a new immutable value

```
long factorial(long n) {  
    return LongStream  
        .rangeClosed(1, n)  
        .parallel()  
        .reduce(1, (a, b) -> a * b,  
                (a, b) -> a * b);  
}
```



`reduce()` combines two immutable values (e.g., long or Long) & produces a new one

Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
 - `reduce()` creates a new immutable value
 - `collect()` mutates an existing value

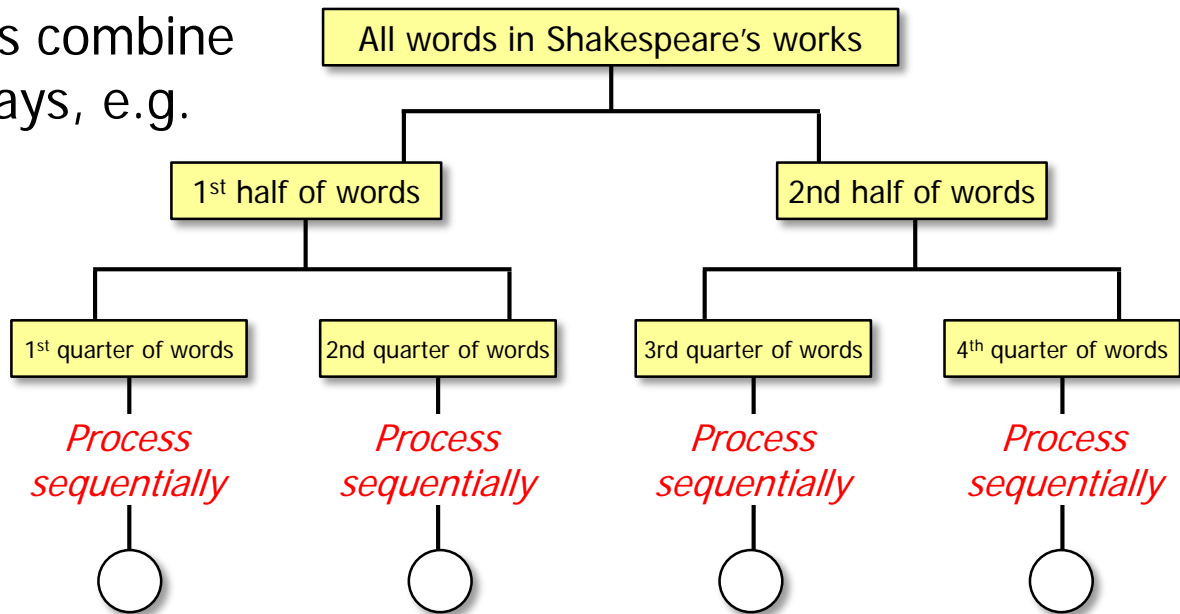


Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.

- `reduce()` creates a new immutable value
- `collect()` mutates an existing value

```
List<CharSequence>  
    uniqueWords =  
        getInput(sSHAKESPEARE),  
            "\\s+")  
        .parallelStream()  
        ...  
        .collect(toCollection(TreeSet::new));
```



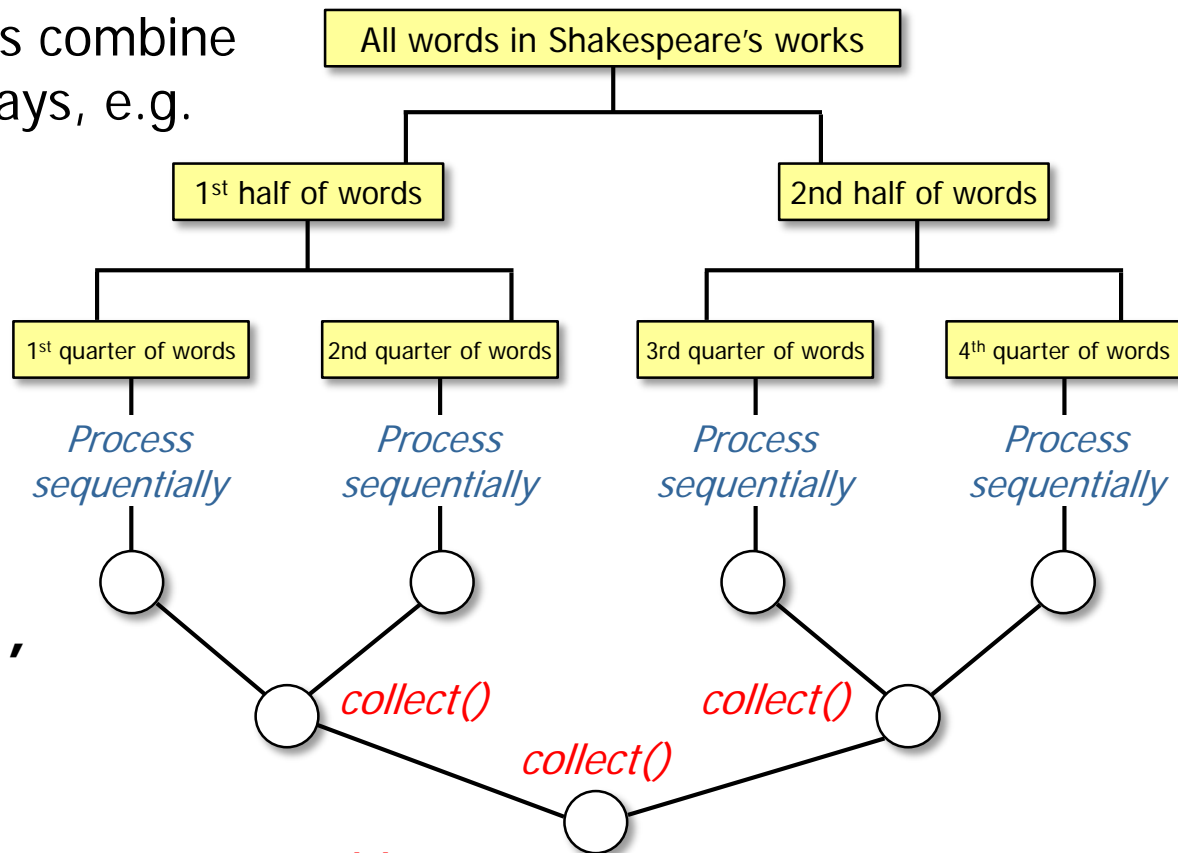
See github.com/douglasraigschmidt/LiveLessons/tree/master/Java8/ex14

Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.

- `reduce()` creates a new immutable value
- `collect()` mutates an existing value

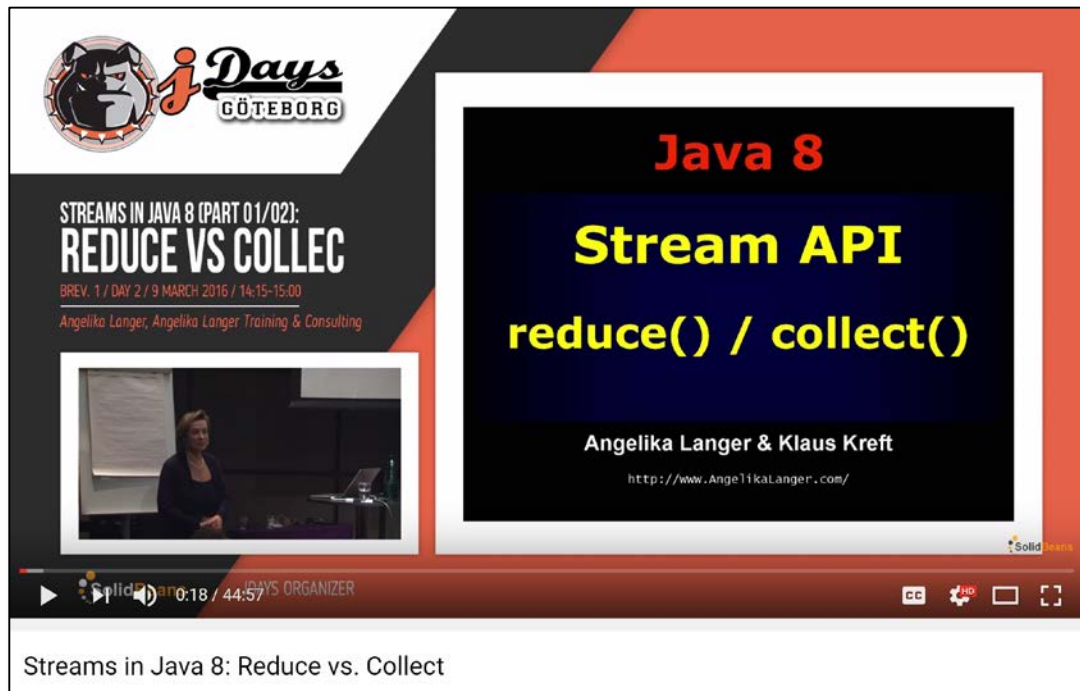
```
List<CharSequence>  
    uniqueWords =  
    getInput(sSHAKESPEARE),  
            "\\s+")  
    .parallelStream()  
    ...  
    .collect(toCollection(TreeSet::new));
```



`collect()` mutates a container to accumulate the result it's producing

Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online



The screenshot shows a video player interface. The main content is a presentation slide with a dark blue background and yellow text. The slide title is "Java 8 Stream API" and the subtitle is "reduce() / collect()". The presenter's name is "Angelika Langer & Klaus Kreft" and the website "http://www.AngelikaLanger.com/" is listed at the bottom. The video player has a red progress bar and a timestamp of 0:18 / 44:57. The video title at the bottom is "Streams in Java 8: Reduce vs. Collect".

**Streams in Java 8 (PART 01/02):
REDUCE VS COLLEC**
BREV. 1 / DAY 2 / 9 MARCH 2016 / 14:15-15:00
Angelika Langer, Angelika Langer Training & Consulting

Java 8
Stream API
reduce() / collect()

Angelika Langer & Klaus Kreft
<http://www.AngelikaLanger.com/>

0:18 / 44:57

Streams in Java 8: Reduce vs. Collect

See www.youtube.com/watch?v=oWIWEKNM5Aw

Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

```
void buggyStreamReduce
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```

Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

This code fails when `parallel()` is used since `reduce()` expects to do an "immutable" reduction

```
void buggyStreamReduce
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```

Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

There are race conditions since there's just a single shared `StringBuilder`, which is not thread-safe..

```
void buggyStreamReduce
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
                StringBuilder::append,
                StringBuilder::append)
        .toString();
}
```

Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

A stream can be dynamically switched to "parallel" mode!

```
void buggyStreamReduce
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```

Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online, e.g.
 - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
- Beware of issues related to association & identity

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
                (x, y) -> x - y);  
}
```

```
void testBuggySequentialSum() {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(1L,  
                Math::addExact);  
}
```

Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online, e.g.
 - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
- Beware of issues related to association & identity

This code fails for a parallel stream since subtraction is not associative

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
            (x, y) -> x - y);  
}
```

```
void testBuggySequentialSum() {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(1L,  
            Math::addExact);  
}
```


Combining Results in a Parallel Stream

- More discussion about `reduce()` vs. `collect()` appears online, e.g.
 - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
- Beware of issues related to association & identity

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
            (x, y) -> x - y);  
}
```

```
void testSum(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
            // Could use (x, y) -> x + y  
            Math::addExact);  
}
```

*This code fails if
identity is not 0L*

The "identity" of an OP is defined as "identity OP value == value"

Combining Results in a Parallel Stream

More discussion about `reduce()` vs. `collect()` appears online, e.g.

- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
- Beware of issues related to association & identity
- Be aware of custom collectors

`mList.stream().collect(Collector.of`

SearchResults's custom collector formats itself

```
(( ) -> new StringJoiner("|"),  
(j, result) -> j.add(result.toString()),  
StringJoiner::merge,  
StringJoiner::toString));
```



See www.youtube.com/watch?v=H7VbRz9aj7c

End of Overview of Java 8 Parallel Streams (Part 2)