

Overview of Java 8 Parallel Streams

(Part 3)

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

Institute for Software
Integrated Systems

Vanderbilt University
Nashville, Tennessee, USA



Learning Objectives in this Part of the Lesson

- Recognize how Java 8 applies aggregate operations & functional programming features in the parallel streams framework
- Be able to avoid concurrency hazards in parallel streams
- Understand how a parallel stream splits its elements recursively, processes them independently & combines the results
- Know when to use parallel streams



Learning Objectives in this Part of the Lesson

- Recognize how Java 8 applies aggregate operations & functional programming features in the parallel streams framework
- Be able to avoid concurrency hazards in parallel streams
- Understand how a parallel stream splits its elements recursively, processes them independently & combines the results
- Know when to use parallel streams
 - & when *not* to use parallel streams



When to Use Java 8 Parallel Streams

When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions



When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent

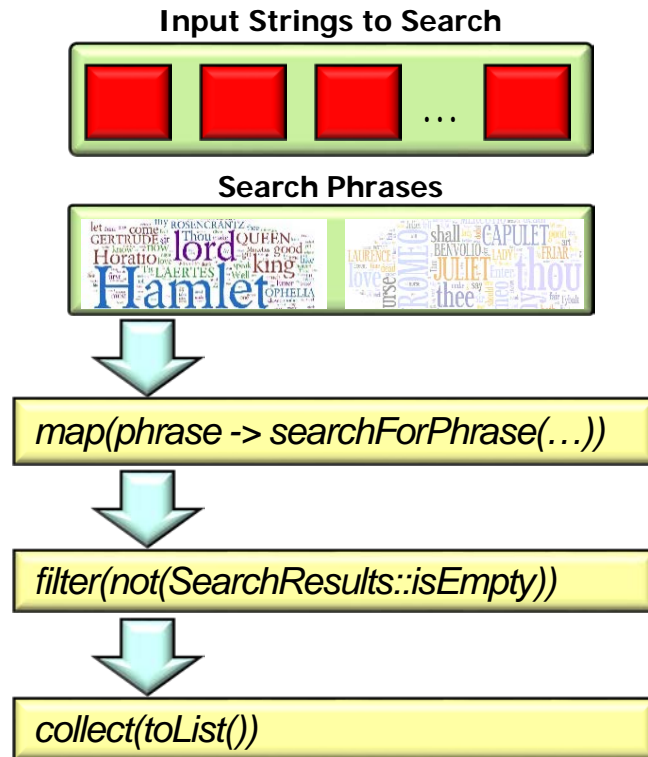
“Embarrassingly parallel” tasks have little/no dependency or need for communication between tasks or for sharing results between them



See en.wikipedia.org/wiki/Embarrassingly_parallel

When to Use Java 8 Parallel Streams

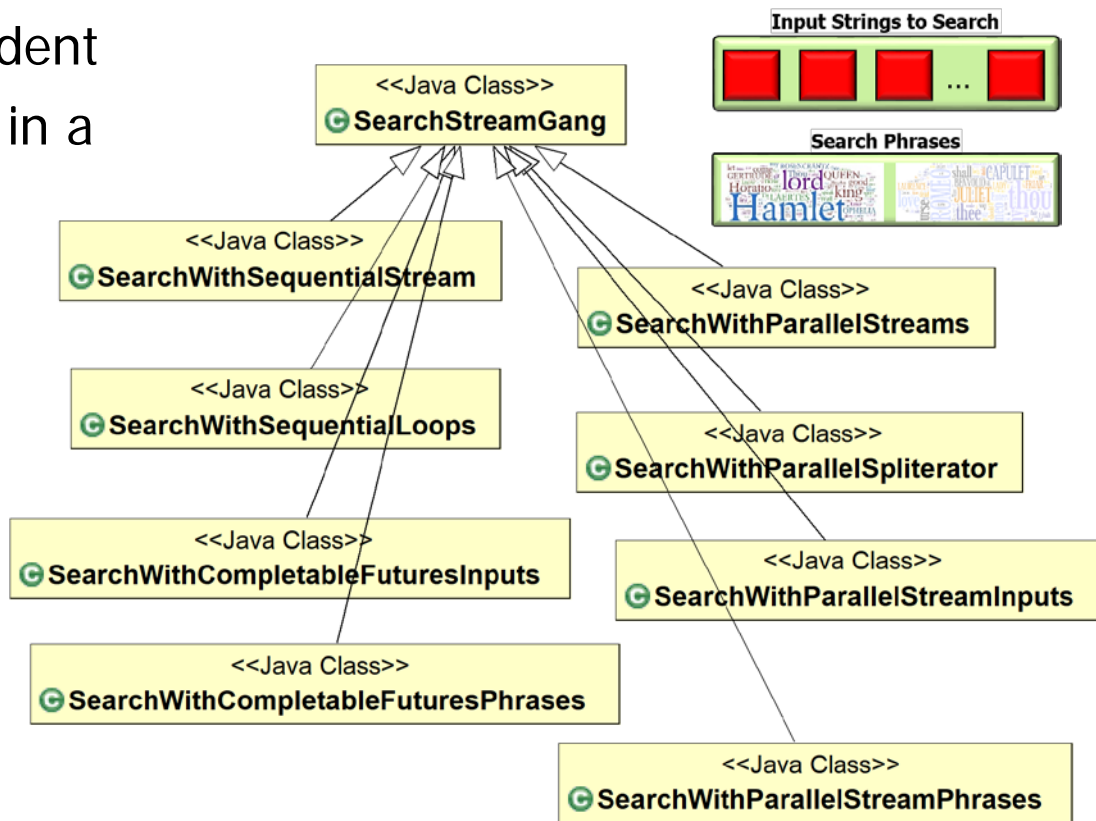
- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - e.g., searching for phrases in a list of input strings



See github.com/douglascraigschmidt/LiveLessons/tree/master/SearchStreamGang

When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - e.g., searching for phrases in a list of input strings
-

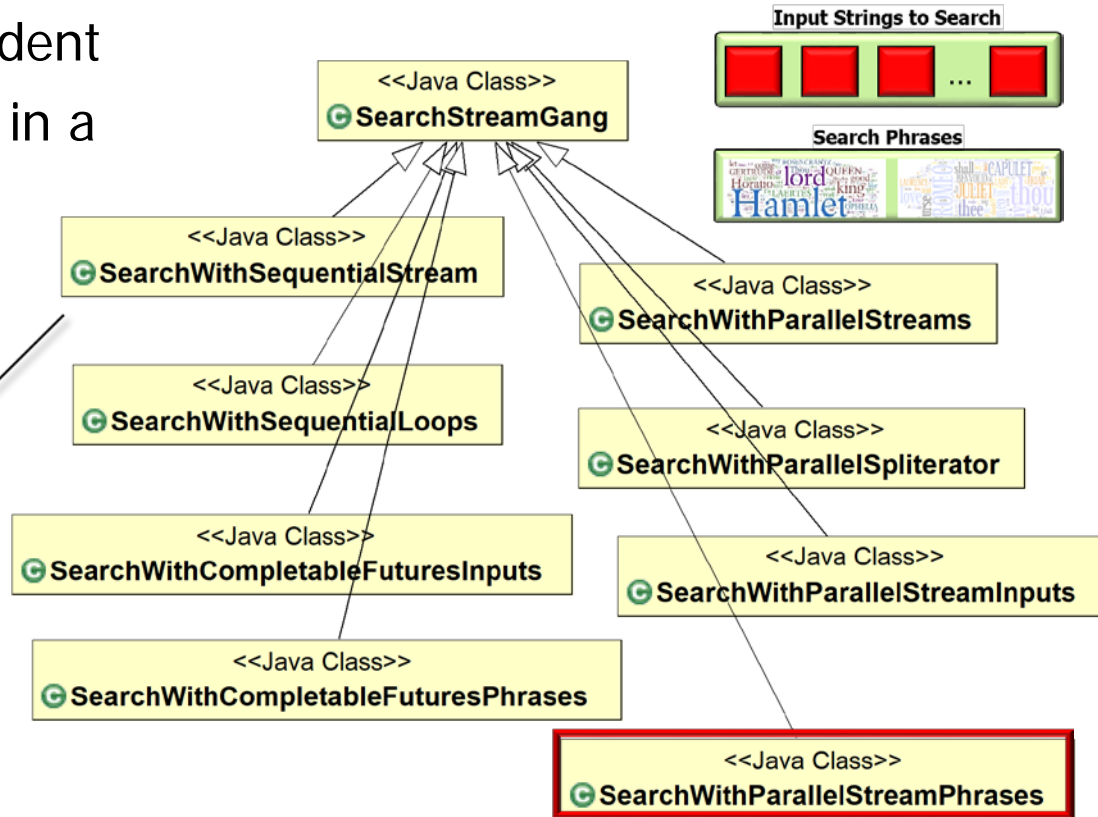


When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - e.g., searching for phrases in a list of input strings

Parallel streams can be used to:

- search chunks of phrases concurrently*
- search chunks of input concurrently*
- search chunks of each input string concurrently*

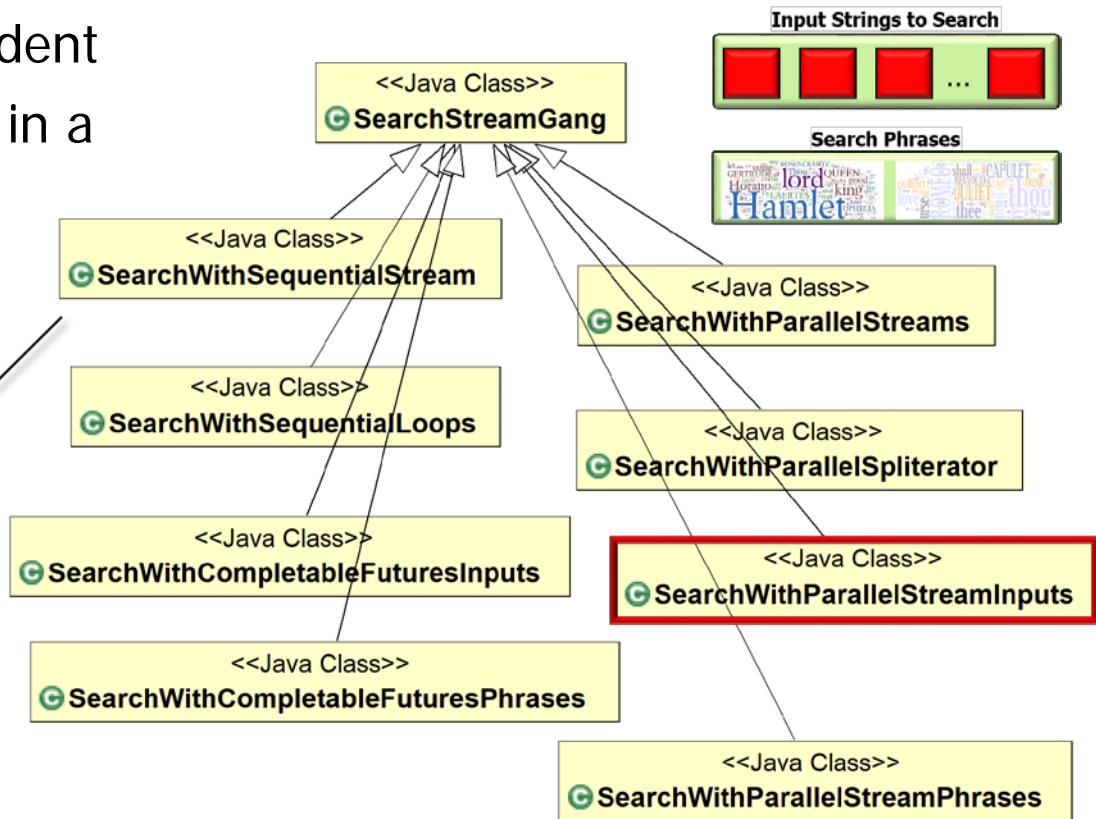


When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - e.g., searching for phrases in a list of input strings

Parallel streams can be used to:

- search chunks of phrases concurrently*
- search chunks of input concurrently*
- search chunks of each input string concurrently*

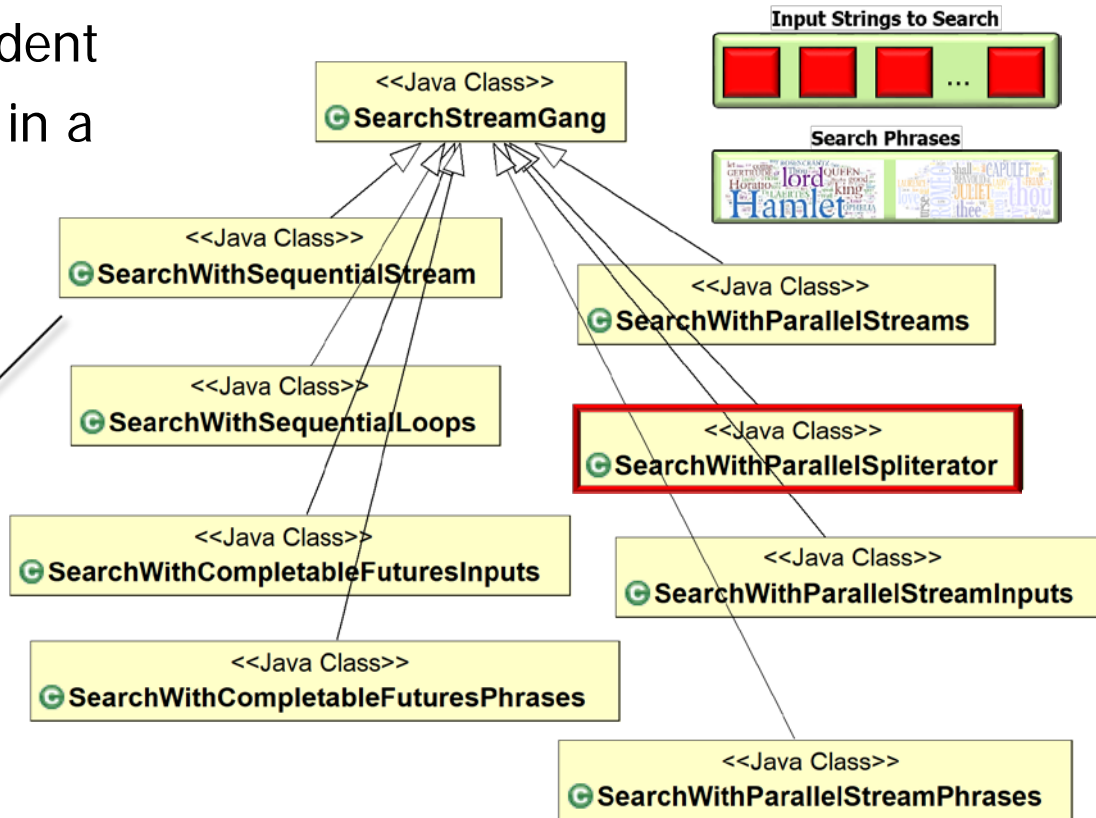


When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - e.g., searching for phrases in a list of input strings
-

Parallel streams can be used to:

- *search chunks of phrases concurrently*
- *search chunks of input concurrently*
- *search chunks of each input string concurrently*

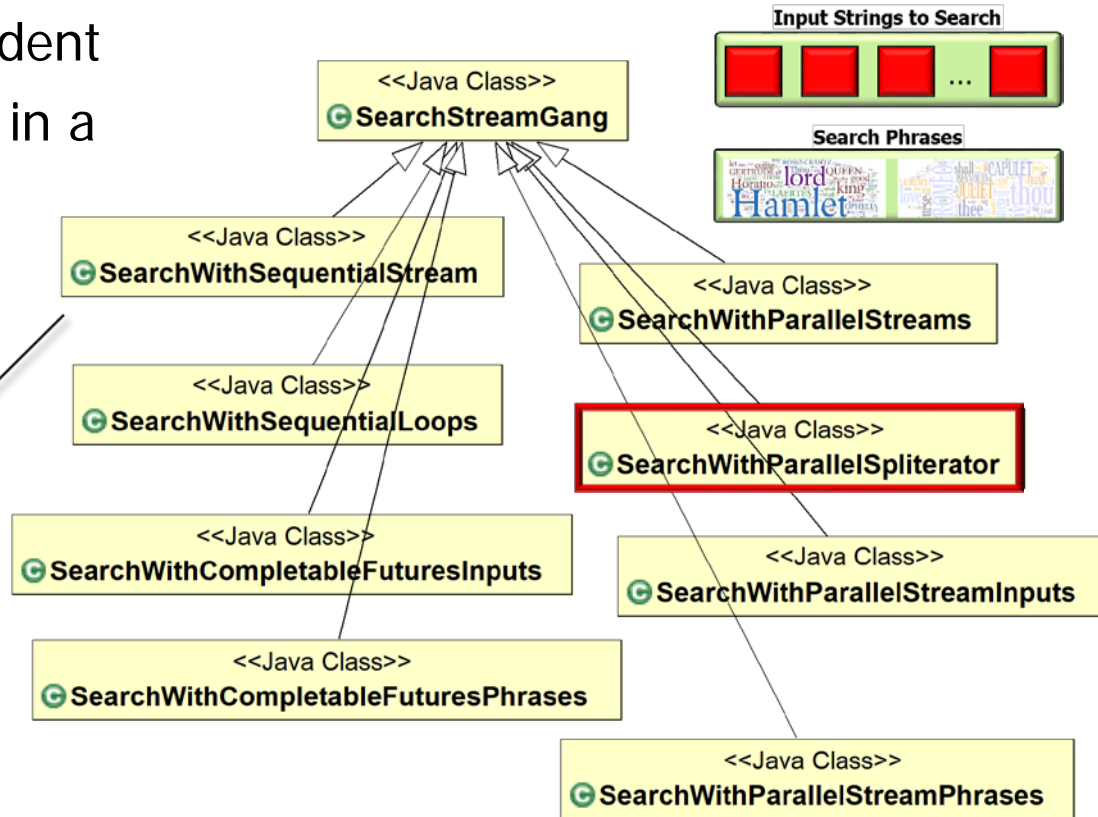


When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - e.g., searching for phrases in a list of input strings

Parallel streams can be used to:

- search chunks of phrases concurrently*
- search chunks of input concurrently*
- search chunks of each input string concurrently*



SearchWithParallelSpliterator is the most aggressively concurrent strategy!

When to Use Java 8 Parallel Streams


- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - When behaviors are expensive computationally and/or applied to many elements of efficiently splittable data structures



See www.ibm.com/developerworks/library/j-java-streams-5-brian-goetz

When to Use Java 8 Parallel Streams


- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - When behaviors are expensive computationally and/or applied to many elements of efficiently splittable data structures
 - i.e., the “NQ” model



- *N is the # of data items to process per thread*
- *Q quantifies how CPU-intensive the processing is*

When to Use Java 8 Parallel Streams

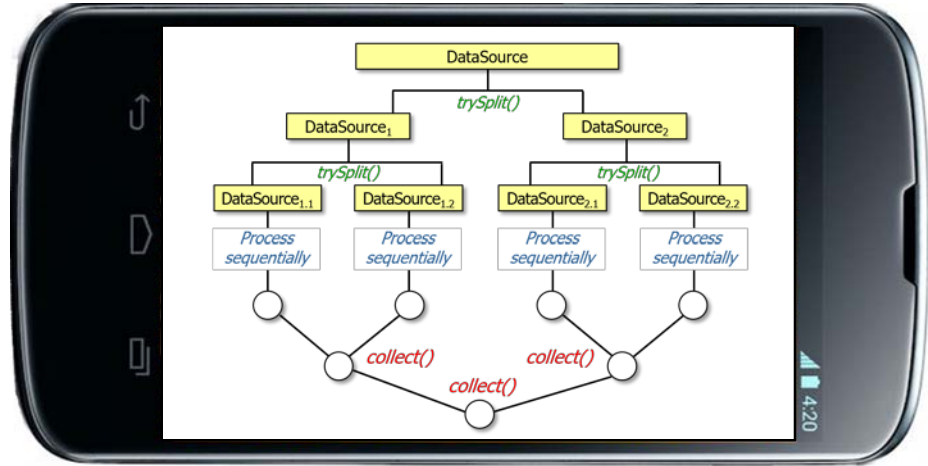
- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - When behaviors are expensive computationally and/or applied to many elements of efficiently splittable data structures
 - i.e., the “NQ” model



e.g., PhraseMatchSpliterator splits input strings into chunks that it searches for regex matches in parallel


When to Use Java 8 Parallel Streams

- Java 8 parallel streams are useful in some (but by no mean all) conditions
 - When behaviors are independent
 - When behaviors are expensive computationally and/or applied to many elements of efficiently splittable data structures
- If there are multiple cores



See blog.oio.de/2016/01/22/parallel-stream-processing-in-java-8-performance-of-sequential-vs-parallel-stream-processing

- If the right conditions apply then Java 8 parallel streams can scale up nicely on multi-core/many-core processors



Ending SearchStreamGangTest

See www.infoq.com/presentations/parallel-java-se-8

When Not to Use Java 8 Parallel Streams

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs



See www.ibm.com/developerworks/library/j-java-streams-5-brian-goetz

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
- The source is expensive to split or splits unevenly

```
List<CharSequence> arrayAllWords =  
    TestDataFactory.getInput  
        (SSHAKESPEARE_WORKS, "\\s+");
```

```
List<CharSequence> listAllWords =  
    new LinkedList<>(arrayAllWords);
```

Make a LinkedList that contains all words in the works of Shakespeare

```
arrayAllWords.parallelStream()  
    .count();
```

```
listAllWords.parallelStream()  
    .count();
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
- The source is expensive to split or splits unevenly

```
List<CharSequence> arrayAllWords =  
    TestDataFactory.getInput  
        (SSHAKESPEARE_WORKS, "\\s+");
```

```
List<CharSequence> listAllWords =  
    new LinkedList<>(arrayAllWords);
```

*The ArrayList parallel stream
is much faster than the
LinkedList parallel stream*

```
arrayAllWords.parallelStream()  
    .count();
```

```
listAllWords.parallelStream()  
    .count();
```

LinkedList splits poorly since finding the midpoint requires traversing ½ the list

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
- The source is expensive to split or splits unevenly

The ArrayList spliterator runs in $O(1)$ constant time

```
class ArraySpliterator {  
    public Spliterator<T> trySplit(){  
        int lo = index, mid =  
            (lo + fence) >>> 1;  
        return (lo >= mid)  
            ? null  
            : new ArraySpliterator<>  
                (array,  
                 lo, index = mid,  
                 characteristics);  
    }  
}
```

See grepcode.com/file/repository.grepcode.com/java/root/jdk/openjdk/8-b132/java/util/Spliterators.java

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
- The source is expensive to split or splits unevenly

The LinkedList spliterator runs in $O(n)$ linear time

```
class LLSpliterator {  
    public Spliterator<E> trySplit(){  
        ...  
        int n = batch + BATCH_UNIT;  
        ...  
        Object[] a = new Object[n];  
        int j = 0;  
        do { a[j++] = p.item; }  
        while ((p = p.next) != null  
                && j < n);  
        ...  
        return Spliterators  
            .spliterator(a, 0, j,  
                Spliterator.ORDERED);  
    }  
}
```

See grepcode.com/file/repository.grepcode.com/java/root/jdk/openjdk/8-b132/java/util/LinkedList.java

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data

```
class ParallelStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n)  
            .parallel() ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    }  
}
```

...

```
class SequentialStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n) ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    }  
}
```

...

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data

The overhead of creating a parallel stream is > than the benefits of parallelism for small values of 'n'

```
class ParallelStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n)  
            .parallel() ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    }  
    ...  
}
```

```
class SequentialStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n) ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    }  
    ...  
}
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data
 - Combining partial results is costly

```
List<CharSequence> allWords =  
    new LinkedList<>  
        (TestDataFactory.getInput  
            (SSHAKESPEARE_DATA_FILE,  
             "\\s+"));  
...  
  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
        ...  
        .collect(toCollection  
            (TreeSet::new));
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data
 - Combining partial results is costly

Performance will be poor due to the overhead of combining partial results for a Set in a parallel stream

```
List<CharSequence> allWords =  
    new LinkedList<>  
        (TestDataFactory.getInput  
            (SSHAKESPEARE_DATA_FILE,  
             "\\s+"));  
...  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
        ...  
        .collect(toCollection  
                    (TreeSet::new));
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data
 - Combining partial results is costly

The combining cost can be alleviated by the amount of work performed per element (i.e., the "NQ model")

```
List<CharSequence> allWords =  
    new LinkedList<>  
        (TestDataFactory.getInput  
            (SSHAKESPEARE_DATA_FILE,  
             "\\s+"));  
  
...  
  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
        ...  
        .collect(toCollection  
            (TreeSet::new));
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data
 - Combining partial results is costly
- A Java 8 feature doesn't enable sufficient exploitable parallelism

```
List<Double> result = Stream
    .iterate(2, i -> i + 1)
    .parallel()
    .filter(this::isEven)
    .limit(number)
    .map(this::findSqrt)
    .collect(toList());
```

```
List<Double> result = LongStream
    .range(2, (number * 2) + 1)
    .parallel()
    .filter(this::isEven)
    .mapToObj(this::findSqrt)
    .collect(toList());
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data
 - Combining partial results is costly
- A Java 8 feature doesn't enable sufficient exploitable parallelism

*Stream.iterate() & limit()
split & parallelize poorly...*

```
List<Double> result = Stream
    .iterate(2, i -> i + 1)
    .parallel()
    .filter(this::isEven)
    .limit(number)
    .map(this::findSqrt)
    .collect(toList());
```

```
List<Double> result = LongStream
    .range(2, (number * 2) + 1)
    .parallel()
    .filter(this::isEven)
    .mapToObj(this::findSqrt)
    .collect(toList());
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data
 - Combining partial results is costly
- A Java 8 feature doesn't enable sufficient exploitable parallelism

LongStream.range() splits nicely & thus runs efficiently in parallel

```
List<Double> result = Stream
    .iterate(2, i -> i + 1)
    .parallel()
    .filter(this::isEven)
    .limit(number)
    .map(this::findSqrt)
    .collect(toList());
```

```
List<Double> result = LongStream
    .range(2, (number * 2) + 1)
    .parallel()
    .filter(this::isEven)
    .mapToObj(this::findSqrt)
    .collect(toList());
```

When Not to Use Java 8 Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly
 - The startup costs of parallelism overwhelm the amount of data
 - Combining partial results is costly
 - A Java 8 feature doesn't enable sufficient exploitable parallelism
 - There aren't many/any cores



Older computing devices just have a single core, which limits available parallelism

End of Overview of Java 8 Parallel Streams (Part 3)